

Worker-Side Discrimination: Beliefs and Preferences

Evidence from an Information Experiment with Jobseekers ^{*†}

Md Moshi Ul Alam [§]

Queen’s University

Mehreen Mookerjee

Zayed University

Sanket Roy

American University of Sharjah

November 2022

Abstract

In this paper, we provide novel evidence on the distribution of workers’ preferences on manager gender and their beliefs on managers’ mentoring ability, which impacts workers’ careers. We design and conduct a within-worker information experiment, embed it in a hypothetical job choice survey for jobseekers and use it to estimate a job choice model with worker preferences and beliefs. In the absence of information on manager quality, on average, workers are indifferent to manager gender. However, upon receiving information on manager mentorship ability, workers prefer to work for female managers—as exhibited by their willingness to receive 1.3–2.2% lower wages on average for the associated jobs. Hence, absent additional information on mentorship skill, workers on average believe that female managers’ mentoring ability is worse than male managers’, with the magnitude of this evaluation corresponding to a wage differential of 1.6% of average annual wages for the associated jobs. These averages mask rich heterogeneity. We find that 60% of workers prefer to work for female managers, and in the absence of information on mentorship ability, 62% believe male managers to be better mentors. An ex post survey directly eliciting worker beliefs corroborates this finding. We find policy-relevant heterogeneity by maternal education level, parental employment status and worker major.

^{*}For guidance, we thank Chris Taber, Jeff Smith, Jesse Gregory and Matt Wiswall. For helpful comments, we thank Peter Arcidiacono, John Kennan, Steve Lehrer, Rasmus Lentz, Corina Mommaerts, Laura Schecter, and Basit Zafar and many seminar and conferece participants. For helpful discussions, we thank Monica Agarwal, Gary Baker, Amrita Kulka, Dennis McWeeny and Elan Segarra. We are grateful to Suraj Das, Iman Kalyan Ghosh, Kabir Rana, Shreya Seal and Aishani Sengupta for excellent research assistance in administering the pilots, ex post debriefs and final survey. Any errors are our own.

[†]We acknowledge funding support for this research by the American University of Sharjah (grant: FRG19-M-B41; IRB approval protocol number: 19-509).

[‡]A previous version of this paper was circulated with the corresponding author’s contact as alam4@wisc.edu and affiliation as the University of Wisconsin – Madison.

[§]Corresponding author: alam.m@queensu.ca

1 Introduction

Managers differ considerably in their ability to manage workers, which directly impacts workers' careers (Frederiksen, Kahn & Lange 2020). Managers with a high ability to manage workers have much lower attrition and turnover among their subordinates (Hoffman & Tadelis (2021), Lazear, Shaw & Stanton (2015))¹. While workers value many nonpecuniary benefits of their jobs (Dey & Flinn (2005), Blau & Kahn (2017), Mas & Pallais (2017), Wiswall & Zafar (2018), Taber & Vejlin (2020)), it remains an open question as to how manager gender and ability directly influences workers' job choice.

The job choice of workers, especially jobseekers, depends on—in addition to their preferences—their beliefs (Robinson (1933), Conlon, Pilossoph, Wiswall & Zafar (2018), Jäger, Roth, Roussille & Schoefer (2021)) because they may not have complete information about their managers. Driven by their preferences and beliefs, should workers prefer not to work for women, they would need to be paid a wage premium to work for female managers. In equilibrium, this could lower the rate at which women are hired or promoted to managerial positions and thus generate a glass ceiling. Thus, worker preferences and beliefs are direct objects of interest given the high turnover costs of replacing and training workers, especially in tight labor markets.

In this paper, we provide novel evidence on the distributions of worker preferences on manager gender and of worker beliefs on managerial ability. We define worker-side discrimination—in the spirit of a compensating differential (Rosen 1986)—as a form of selection, where individuals are willing to forgo wages to work for their preferred managers in an otherwise-identical job.² This willingness to trade off wages is driven by their preferences on observable attributes and beliefs on unobservable attributes that they care about but do not have information on.

To identify the distribution of preferences, we follow the literature to design and conduct a hypothetical job choice survey to ensure that demand-side selection, labor market frictions and other omitted variables in general do not confound our results (Blass, Lach & Manski (2010), Wiswall & Zafar (2018), Ameriks, Briggs, Caplin, Shapiro & Tonetti (2020), Fuster, Kaplan & Zafar (2021), Koşar, Ransom & Van der Klaauw (2021), Koşar, Şahin & Zafar (2021)).³ Hypothetical choice methods are attractive because they can allow unre-

¹Manager heterogeneity also directly impacts workers' wages and wage inequality within firms (He & le Maire (2020), Acemoglu et al. (2022)).

²Becker (1971) conceptualized worker discrimination in the form of worker disutility from working for a specific group of employers. We extend the concept to incorporate worker beliefs and a tangible measure using compensating differentials in wages.

³Different workers may have different preferences on various dimensions of job attributes, many of which are unobservable to the researcher. Such preferences are very difficult to isolate using data on real-

stricted forms of preference heterogeneity (Blass, Lach & Manski 2010) while being able to hold fixed attributes not considered in the survey through instructions (Wiswall & Zafar (2018), Koşar, Ransom & Van der Klaauw (2021), Koşar, Şahin & Zafar (2021)) and to document strong mapping between stated and actual choices (Mas & Pallais (2017), Wiswall & Zafar (2018), Parker & Souleles (2019)).

To identify the distribution of beliefs, we embed a novel within-worker information experiment where we exogenously vary the observability of managerial ability. We define managerial ability as the manager’s mentorship quality, motivated by a vast literature providing consistent evidence that mentorship has substantial positive impacts on human capital accumulation (Falk, Kosse & Pinger 2020), wage expectations (Boneva, Buser, Falk, Kosse et al. 2021), productivity (Blau, Currie, Croson & Ginther 2010), promotions (Lyle & Smith 2014), and the workforce composition of firms and can help minorities break through glass ceilings (Athey, Avery & Zemsky (2000), Müller-Itten & Öry (2022)).⁴ We quantify mentorship of a manager as a rating on a five-point scale, motivated by a recent trend of rating managers.⁵

We conduct this hypothetical choice survey and the information experiment among jobseeking students enrolled at a highly selective university who are one year away from graduating. We present respondents with twenty hypothetical job choice scenarios sequentially. In each choice scenario, we ask respondents to choose one out of three jobs. We exogenously vary these jobs along realistic attributes (annual wages, flexible hours, manager gender and manager mentorship quality) and cover the support of these attributes over the twenty different job choice scenarios.⁶ In each scenario, respondents are asked to choose their most preferred job and then report the compensating differential in wages—a nonparametric cardinal measure—that would make them indifferent between the preferred job and the other jobs.

Extracting beliefs on manager quality through direct elicitation could be difficult, especially for realized job choices. However, data on realized job choices do have their own advantages, especially for helping us understand employer discrimination. This is because employers on average care about a consistent set of attributes in their workers.

⁴In the same spirit of in-group mentoring, through the American Economic Association’s (AEA’s) official mentoring program CeMENT, senior women faculty mentor junior women faculty. See the AEA’s dedicated webpage on mentoring [here](#).

⁵Many firms such as Google, e-Bay, and Amazon collect anonymous surveys from employees where they are asked to rate their managers. Comparably, Completed, TheJobCrowd and Kunukunu are some of notable start-ups that provide manager ratings analogous to Glassdoor’s firm ratings.

⁶Conceptually, each hypothetical scenario could be thought of as a market. Choice in a market provides individual demand in that market. Survey data on choices over multiple scenarios varying attributes over their support allow us to trace out the individual demand curve. This generates panel data on choices and compensating differentials over the support of the job attributes, which provide the identifying variation to estimate highly flexible models.

pecially if we are concerned that the responses may be affected by social desirability bias. Hence, within the hypothetical job choice scenarios, we introduce a novel within-individual information experiment where we exogenously vary the observability of manager mentorship skill. The information experiment works as follows. In the first ten job choice scenarios, every individual observes three jobs in each scenario with different attributes: annual wages, flexibility of hours and manager name. As an attribute, mentorship is mentioned, but the data are shown to be unavailable. We call these first ten scenarios "incomplete scenarios" throughout the rest of the paper, given that the mentorship rating is not observable. In the last ten scenarios, individuals observe jobs with all of the above attributes as well as a manager mentorship rating. We call these last ten scenarios "complete scenarios". We elaborate on the key highlights of our design later in the paper.

We use our unique panel data on choices and compensating differentials to estimate a structural model of job choice where we estimate worker preference and belief parameters in monetary value—as a willingness to forgo wages. Identification is achieved as follows. Each worker forms expected utilities while choosing a job. In the incomplete scenarios, workers implicitly form expectations on the mentorship rating because they do not observe it. Thus, their responses are a function of both their preferences and their beliefs on mentorship conditional on other attributes. In contrast, in the complete scenarios, since individuals observe all attributes, their responses are a function of only their preferences. We instruct respondents in every scenario that the jobs do not vary in attributes not mentioned in the survey (Wiswall & Zafar (2018), Koşar, Şahin & Zafar (2021), Koşar, Ransom & Van der Klaauw (2021)) and that the reported compensating differentials only increase wages without changing anything else about the job.⁷ Thus, the variation in compensating differentials within the complete scenarios identifies preferences. The variation in compensating differentials between the complete and incomplete scenarios resulting from the information experiment then isolates beliefs from preferences. Finally, with the preference and belief parameters identified for each worker, we can identify the corresponding distributions.

We find that in the absence of information on manager mentorship, such that choices and compensating differentials are driven by *both preferences and beliefs*, workers are indifferent between male and female managers. However, with information on manager mentorship skill, such that choices and compensating differentials are driven by *only pref-*

⁷This is one of the key advantages of using the hypothetical choice methodology over audit study field experiments. We also incorporate direct and indirect questions later in the survey to test how closely these instructions are followed.

erences, workers prefer to work for female managers. On average, workers are willing to give up 1.7% of their average annual wages to work for female managers. Hence, in the absence of information on manager quality, workers believe female managers to be worse mentors. We estimate the value of these negative beliefs on female managers' mentorship ability to be equivalent on average to 1.6% of workers' average annual wages.

An important finding from our within-worker information experiment is that there exists rich heterogeneity in the underlying distribution of worker preferences and beliefs on managers' gender and mentoring ability. Approximately 62% of individuals prefer to work for female managers. Approximately 60% of individuals believe female managers to be worse mentors than male managers in the absence of information on mentorship skill. Individuals majoring in engineering are more likely to prefer to work for female managers than those majoring in the humanities. Individuals whose mothers are weakly more educated than their fathers are less likely to have negative beliefs about female managers. We also find heterogeneity by the joint employment status of the individual's parents. Unlike [Flory, Leibbrandt & List \(2015\)](#) and [Wiswall & Zafar \(2018\)](#), we do not find evidence of differences in average preferences and beliefs by respondent gender.

In addition to our information experiment within the hypothetical choice survey, we collect further data to support our results. After the hypothetical choice scenarios, we ask questions that directly elicit respondents' beliefs. We ask respondents to report their expected mentorship rating of managers in ten hypothetical jobs while we exogenously vary manager names, flexibility of hours and annual wages. This allows us to corroborate the results on beliefs on manager mentorship from the information experiment in the job choice survey. Here, too, we find evidence of average negative beliefs regarding female manager mentorship skill similar to what we found in the information experiment involving the twenty job scenarios. After going through all incomplete and complete choice scenarios, on average, respondents still report negative beliefs on female manager mentorship skill when asked directly. This corroborating result tells us that the individual responses in our information treatment indirectly eliciting their beliefs are potentially robust to social desirability bias.

Our paper contributes to multiple strands of the literature. First, to the best of our knowledge, this is the first paper to provide evidence on the distribution of worker preferences and beliefs on manager gender and mentorship.⁸ While [Flory, Leibbrandt & List](#)

⁸Recent work by [Abel \(2019\)](#), [Abel & Buchman \(2020\)](#) and [\(Ayalew, Manian & Sheth 2021\)](#) focuses on how manager's achievements in their own jobs impacts how likely workers are to follow their advice. Unlike the focus of this literature on managers' ability in their own job, our focus is on managers' mentorship quality, in view of the evidence on the significant impacts of managers' mentoring ability on their subordinates' labor market outcomes ([Hoffman & Tadelis 2021](#)) and the evidence for the Peter principle ([Benson](#),

(2015) find no evidence of a role of manager gender in application decisions, our analysis reveals that this finding is sensitive to the information that workers have about managers and that there exists substantial underlying heterogeneity in this regard. Second, the literature on discrimination usually deals with average discrimination driven by beliefs (statistical and biased beliefs) and by preferences (taste-based) separately (Charles & Guryan (2008), Guryan & Charles (2013), Lang & Lehmann (2012), Bertrand & Duflo (2017)). Kline, Rose & Walters (2021) estimate the distribution of racial discrimination, but they consider discrimination by firms toward workers. To the best of our knowledge, this paper is the first to explicitly allow for discrimination driven by both worker beliefs and preferences and to estimate their distributions. Our design generating unique panel data on compensating differentials allows us not only to test for statistical discrimination (Altonji & Pierret (2001), Lange (2007), Agan & Starr (2018)) but also to quantify it (Bohren, Imas & Rosenberg 2019) as a measure of the willingness to forgo wages.⁹

Our next contribution is methodological. The literature using the stated-preference methodology estimates preference parameters in scenarios where individuals have information on all attributes of interest while other attributes are held fixed with respect to various objects of choice: e.g., electricity services (Blass, Lach & Manski 2010), jobs (Wiswall & Zafar 2018), residential locations (Koşar, Ransom & Van der Klaauw 2021), political candidates (Delavande & Manski 2015), and insurance products (Boyer, De Donder, Fluet, Leroux & Michaud 2017). On this aspect, our paper is closest to Wiswall & Zafar (2018). We differ from this work by using an information experiment and eliciting wage compensating differentials between jobs instead of choice probabilities for each job. Given our design, these compensating differentials allow us to jointly estimate the complete and incomplete scenarios to recover the distribution of preferences and beliefs without making additional assumptions on the variance of the error terms of the job choice model. We can do this because the value of a dollar remains a dollar irrespective of whether the scenario is complete or incomplete.¹⁰ This allows us to directly estimate and interpret the preference and belief parameters as measures of willingness to forgo wages. Later in the paper, we elaborate on the similarities and the differences between asking

Li & Shue 2019)—namely, that high-ability workers, upon being promoted, do not necessarily become managers with high managerial ability.

⁹Bohren, Imas & Rosenberg (2019) additionally distinguish between discrimination resulting from correct beliefs and that resulting from incorrect beliefs in their experimental set-up studying the evolution of discrimination. We cannot take this route because we do not have access to data on the population distribution of mentorship quality that could provide the benchmark to test the hypothesis of biased beliefs.

¹⁰If we had asked for choice probabilities, we would need to additionally estimate the relative variances of the error terms in the complete and incomplete scenarios, which could be identified from the relative variances of the choice probabilities in the complete and incomplete scenarios.

for compensating differentials and asking for choice probabilities, focusing on their intuitions, roles in identification, and manner of linking to the instructions provided to the respondents. Third, our question of interest involves identification of beliefs, similar to the setting of [Adams-Prassl & Andrew \(2019\)](#).¹¹ However, we differ by indirectly eliciting beliefs with our information experiment by using incomplete scenarios, thereby providing a new method for belief elicitation and estimation in settings where the researcher may worry that individuals may not report truthfully because of social desirability bias. Additionally, we conduct an ex post survey that directly elicits beliefs to compare our results from the information experiment. We find that while the results are qualitatively similar, male respondents shade down their negative beliefs on female manager mentorship skill when asked directly about their beliefs, relative to the beliefs elicited under our information experiment.

Another novel contribution of our paper is to quantify the demand for manager mentorship. We have consistent evidence on the positive impacts of mentorship on outcomes of mentees in academia, the corporate sector, the military and high schools ([Athey, Avery & Zemsky \(2000\)](#), [Blau, Currie, Croson & Ginther \(2010\)](#), [Lyle & Smith \(2014\)](#), [Falk, Kosse & Pinger \(2020\)](#), [Müller-Itten & Öry \(2022\)](#), [Boneva, Buser, Falk, Kosse et al. \(2021\)](#)). To the best of our knowledge, our paper provides the first estimates of the demand for high-quality mentors in terms of the wages that jobseekers are willing to forgo to work for managers who are better mentors. We estimate that individuals are willing to forgo up to 5.65% of average annual wages for a one-standard-deviation increase in mentorship rating. This result is crucial to interpreting the beliefs on mentorship skill when it is unobservable to workers. If workers did not care about mentorship, then any belief distribution could rationalize the data, resulting in beliefs being fundamentally unidentified. Our job choice model incorporates this feature.

Finally, our work is also a part of the growing literature featuring online surveys and experiments ([Stantcheva 2022](#)) with information treatments for the study of beliefs ([Wiswall & Zafar \(2015\)](#), [Kuziemko, Norton, Saez & Stantcheva \(2015\)](#), [Bordalo, Coffman, Gennaioli & Shleifer \(2016\)](#), [Alesina, Miano & Stantcheva \(2019\)](#), [Boneva & Rauh \(2018\)](#), [Alesina & Stantcheva \(2020\)](#), [Stantcheva \(2021\)](#), [Alesina, Ferroni & Stantcheva \(2021\)](#)).

The paper is organized as follows: Section 2 provides details on the hypothetical job choice survey and the information experiment and highlights the important features of the design. Section 3 describes the sample and raw patterns in the data. Section 4 describes a job choice model of how workers' preferences and beliefs drive their choices and

¹¹Beliefs in [Adams-Prassl & Andrew \(2019\)](#) are probabilities that individuals place on their own future outcomes, whereas in our paper, beliefs are workers' perceptions of potential managers' mentoring ability.

compensating differentials. Section 5 shows identification. Section 6 discusses estimation details and results. Section 7 presents the empirical distribution of beliefs and preferences and shows evidence on their underlying heterogeneity. Section 8 discusses the validity of the estimates of the belief parameters and further robustness checks. Section 9 discusses potential avenues for future research. Section 10 concludes.

2 Hypothetical Job Choice Survey and Information Experiment

Our hypothetical job choice survey included the following sections in order (1) instructions to the respondents, (2) twenty hypothetical job choice and compensating differential *scenarios* within which the information experiment was embedded, (3) direct belief elicitation and (4) demographic questions. The structure is schematically represented in Figure 1. Below, we describe the design and purpose of each section in detail.

2.1 Instructions

The first part of the survey included definitions of the attributes of jobs shown to individuals in each job as shown in Figure 2 followed by survey instructions as shown in Figure 3. There were two key instructions. First, individuals were to assume that the jobs do not vary on any attribute not mentioned in the survey. Second, when asked to report the minimum increase in annual wages in unchosen jobs required to make them indifferent to their chosen job, individuals were to assume that this increase in wages would not change anything else about the job. After the instructions, individuals were shown two example scenarios to familiarize them with the set-up before they started the main survey. We later reemphasized the instructions within each scenarios, as well.

2.2 Job choice scenarios with compensating differentials

Each *scenario* had two questions: a choice to be made among three hypothetical jobs, followed by a question on the compensating differentials that would make respondents indifferent between the jobs. Twenty such scenarios were administered in the survey to plausibly cover the support of job attributes, with the first 10 being the incomplete scenarios and the next 10 the complete scenarios.¹² Examples of job choice and compensating

¹²Ideally, we would have varied the job attributes along the full range of attributes, but this would have required asking a large number of questions, but only at a large cognitive cost to the respondents. Hence,

differential questions in both an incomplete scenario and a complete scenario are shown in Tables 1 and 2, respectively. Examples adapted to corresponding representative jobs in the USA are also shown in Tables 3 and 4.

The jobs differed in the following attributes, which we exogenously varied: annual wages, availability of flexible hours, manager name and manager mentorship rating. The manager’s mentorship rating—quantified on a five-point scale—was described to the respondents as the average rating provided by the manager’s current workers in an anonymous survey.¹³ We embedded the information experiment as follows: for every respondent, although the mentorship rating was mentioned as an attribute in all 20 scenarios, the first 10 scenarios (the incomplete scenarios) did not have rating data on the manager’s mentorship ability, while the last 10 scenarios (the complete scenarios) did.

Note that for the incomplete scenarios, the rating variable was mentioned but there were no ratings available for the managers. We had to ensure with regard to the wording that we neither primed individuals to think that the rating was indeed different across managers nor made them assume that the rating was the same across all managers. To achieve this, we used the following wording within each incomplete scenario: *“The rating of each manager may be different, but the data are not available.”* We reemphasize that the jobs did not differ on any other attribute not mentioned in the scenarios.

In every scenario, after a job was chosen, the following question asked the respondents to report the compensating wage differentials for the jobs not chosen. For each unchosen job, respondents were asked to specify the minimum increase in annual wages that they would need to choose that job instead. These data provide us with the compensating wage differentials that would make the respondents indifferent between jobs. Individuals could report these on a slider scale that ranged between 0 and 2 lakhs INR (\approx 0 USD to 2857 USD). Individuals were told that if they needed more than 2 lakhs, they could max out the slider and another page would automatically appear asking them how much more they would need.¹⁴

In the 20 scenarios, there were 60 jobs, with half male and half female managers evenly distributed across the complete and the incomplete scenarios. Approximately half of the

to strike a balance between cognitive load and level of variation in job attributes, we chose to administer 20 scenarios. This choice was made after we observed the duration to completion and time spent on each question, especially the latter ones, in our pilots, which differed in the number of scenarios. Wiswall & Zafar (2018) have 16 scenarios. Koşar, Ransom & Van der Klaauw (2021) have three sets of 8, 16 and 24 scenarios.

¹³Anonymity in surveys is standard practice in the employee survey designs used on Amazon, Google and eBay. We thank Will Dobbie for pointing this out. Cai & Wang (2022) also use anonymous employee surveys in their field experiment to communicate employee feedback to treated teams’ managers.

¹⁴We did not find any evidence of any design-induced bunching of reported compensating differentials at the boundary of the slider at 2 lakhs INR (\approx 2857 USD). We thank Jeff Smith for this comment.

jobs had flexible hours, and the other half did not. The average annual wages were 7 lakh INR (\approx \$ 39,444 in PPP),¹⁵ and the average rating of managers in the complete scenarios was 3.41. We varied the attributes subject to the restriction that no job in each scenario was strictly dominant. Table 5 shows summary statistics of the attributes shown over the 60 jobs, and Table 6 shows the balance of attributes between male and female managers.

2.3 Direct belief elicitation

After the information experiment involving the 10 incomplete and 10 complete hypothetical job choice and compensating differentials scenarios, we also directly elicited beliefs on manager mentorship ability. We designed this component to allow us to compare our results on beliefs identified from the information experiment described above with those elicited by asking individuals directly. In this section, individuals were presented with 10 jobs with the manager’s name, annual wages and availability of flexible hours. Individuals were asked to report on a zero-to-five sliding scale the manager mentorship ratings that they expected to be associated with each job, as shown in Figure 7.

2.4 Demographic questions

The final section of our survey asked the respondents demographic questions on their area of study (arts, science or engineering), family income, parental education and occupation. Then, we asked questions specifically designed to allow us to infer whether they had followed our instructions. The survey ended with the choice of mode of online payment for completing the survey.

2.5 Key highlights of the design

In this section, we highlight some of the important aspects of the design of the hypothetical job choice survey and the way in which the embedded information experiment enables us to identify beliefs and preferences from the reported choices and compensating differentials, given the instructions.

By construction, the ability to observe the entire choice set is one of the key advantages of our design. We observe which jobs are chosen and which jobs are not. Additionally, we observe the compensating differentials that make individuals indifferent across all

¹⁵This was the average annual wages of jobs offered to past graduating cohorts of the university attended by the students in the sample. The variance in wages in the jobs that we showed was not particularly high. This mitigates any concerns that some jobs could be interpreted as entry-level and some as senior-level jobs. The attributes were varied such that no job in any scenario turned out to be strictly dominant.

choices in the choice set. This gives us a nonparametric cardinal measure of utilities and thus allows us to avoid making any distributional assumptions on the preference or belief parameters.

The data on the compensating differentials that make individuals indifferent between jobs allow us to directly estimate and interpret the parameters as measures of willingness to pay or to forgo wages.¹⁶

The information treatment is given to every individual. For each individual, we observe the sequence of choices made and the compensating differentials reported over the incomplete scenarios and then over the complete scenarios. This allows us to uncover the distributions of preferences and beliefs and not just the first moment, which would have been the case had we provided the information treatment to a randomly chosen treatment group.

We collect data over twenty scenarios. We do this to cover as much of the support of the job attributes as feasibly possible. Our concerns over the potentially high cognitive load associated with making choices among a large number of options led us to conclude that it was infeasible to ask individuals to choose among a large number jobs within each scenario. Hence, we did this over a panel of scenarios, making them choose and provide compensating differentials over three jobs per scenario. This generated panel data on choices and compensating differentials over jobs that exogenously vary in attributes.

The names of managers used in the survey were common names directly indicative of gender. We used managers' first names only. In the Indian context, last names can reveal caste and religion. Since we wanted to vary only the gender dimension, we did not show any last names, circumventing any potential concerns over differences in perceived gender roles across social classes. This setting is unlike the US context, where first names can be associated with both a race and a gender (e.g., [Bertrand & Mullainathan \(2004\)](#), [Kline, Rose & Walters \(2021\)](#)).

3 Data

We collected data in the second week of April 2020, using an online survey¹⁷ administered to students of a highly selective public university in India. Students eligible to participate

¹⁶This is also possible with data on choice probabilities but requires an additional step to transform the estimates into willingness-to-pay measures. See [Wiswall & Zafar \(2018\)](#).

¹⁷Access to the internet was not a concern for our sample of students studying in a premier university in one of the largest metropolitan cities of India. However, we paid special attention in designing the survey to ensuring that it was mobile-friendly, in consideration of the fact that a small but significant proportion of the target population might not have access to a computer.

in the survey were only those at most one year away from graduating.¹⁸ Upon completion of the survey, participants were paid INR 500 (\approx \$24 in PPP¹⁹) through their preferred online payment mode. See Appendix A.5 for further details on the administration and implementation of the survey.

3.1 Sample selection and description

The total number of participants in our survey was 604, among which 591 completed the survey. Out of these, we dropped 11 respondents who could not be verified as students or completed the survey in less than 15 minutes or both.²⁰ The median time to survey completion was 51.37 minutes. This brought our final sample size to 580.

Table 7 reports sample descriptives. A total of 41.72% of our sample consisted of female students, and the remaining 58.28% were male students.²¹ With respect to majors, 44% were enrolled in a department in the arts faculty, 33.28% in engineering, and the remaining 22.07% in science. Female students were predominantly arts students (67%). Men were predominantly engineering (49%) and science (33%) students.

3.2 Patterns in the raw data

In this subsection, we explore how individuals' choices and reported compensating differentials varied across the complete and incomplete scenarios between jobs with male and female managers. This section is an important precedent to the section where we structure our data to rationalize them in a model involving preferences and beliefs.

Table 8 reports the percentage of jobs chosen by manager gender in both the complete and incomplete scenarios. The first observation is that between male and female respondents, the percentages of jobs chosen with managers of the two genders do not differ substantively. Furthermore, in the absence of information on mentorship ability (in the incomplete scenarios), the percentage of jobs with female managers chosen was not that different from that of jobs with male managers. However, we observe that in the complete scenarios, upon revelation of the manager mentorship information, the percent-

¹⁸In 2020, the total number of enrolled students in the university was 11,064. Out of these, 6,283 were enrolled in undergraduate programs, 2,588 in master's programs, and the remaining 1,193 in MPhil and PhD programs.

¹⁹The purchasing power parity of 1 USD in 2019 is equivalent to 21.07 INR. Source: <https://data.oecd.org/conversion/purchasing-power-parities-ppp.htm>

²⁰The first percentile of duration to survey completion was at 13.89 minutes.

²¹In the survey, we asked individuals their biological sex. We did not ask about gender identifications not coinciding with biological sex.

age of jobs with female managers chosen is 61.1%, which is approximately 20 percentage points higher than that of jobs with male managers.

Table 9 reports the average compensating differentials reported for unchosen jobs with male and female managers and the difference between them along with the associated standard errors and standardized differences. The table reports these numbers separately for the complete and incomplete scenarios. We observe that in the absence of information on manager mentorship skill, individuals on average report compensating differentials required to choose jobs with female managers over jobs with male managers that are higher by 6.3 thousand INR ($\approx \$ 300$). However, this result flips in the complete scenarios. When provided information on the mentorship rating, individuals on average demand 6.1 thousand INR ($\approx \$ 290$) more for unchosen jobs with male managers. Both differences are statistically significant at the 99% level. We should maintain caution in interpreting these numbers because they compare compensating differentials among the set of jobs not chosen. Nevertheless, these numbers provide useful information on the set of unchosen jobs. A more informative way to understand the compensating differential data would be to incorporate the choices made (the extensive margin) and the compensating differentials for the jobs not chosen (the intensive margin) conditional on job attributes. This is what we do in the job choice model.

Before we delve into the model, we provide evidence for a natural question that arises in these contexts on in-group preferences. In particular, are female respondents more likely to choose jobs with female managers?

3.3 Testing for in-group preferences

In this section, we discuss whether women are more likely than men to choose jobs with female managers. This can be answered with data from the complete scenarios using a simple difference-in-differences estimation strategy. We do not include the incomplete scenario data here because we do not want to deal with the omitted variable bias that would arise from how individuals form beliefs on mentorship skill, which they do not observe in the incomplete scenarios.

Individuals are indexed by $i = 1, \dots, N$ and jobs by $j = 1, \dots, J$. Define $Choice_{ijs}$ as an indicator variable which takes value 1 if individual i in scenario s chooses job j and 0

otherwise.

$$\begin{aligned}
Choice_{ijs} = & \delta_0 + \delta_1 \underbrace{\mathbb{I}(g_i = f)}_{\text{female worker}} \underbrace{\mathbb{I}(MG_{j(s)} = f)}_{\text{female manager}} + \delta_2 \underbrace{\mathbb{I}(g_i = f)}_{\text{female worker}} + \delta_3 \underbrace{\mathbb{I}(MG_{j(s)} = f)}_{\text{female manager}} \\
& + Attributes'_{j(s)} \gamma_1 + Demographics'_i \gamma_2 + \lambda_s + e_{ijs}
\end{aligned} \tag{1}$$

Respondent i 's gender is denoted by g_i and the gender of the manager in job j of scenario s by $MG_{j(s)}$. $Attributes_{j(s)}$ is the vector of job attributes associated with job j in scenario s other than manager gender, i.e., annual wages, flexibility of hours and mentorship skill of the manager. $Demographics_i$ is a vector of the individual-level demographics described in the previous section. Our specification includes scenario fixed effects λ_s to leverage the variation in choices made within scenarios resulting from the variation in attributes between jobs within each scenario. We estimate this difference-in-differences equation with a logit model and bootstrap the standard errors at the individual level.

Table 10 shows the marginal effect estimates from equation (1). We find no evidence that female workers are more likely than male workers to choose jobs with female managers. Note that we see this in the raw data as well, where we find no difference in choices or in compensating differentials across male and female respondents in the complete (or incomplete) scenarios. As one would expect, higher wages, availability of flexible hours and better mentorship are associated with a higher likelihood of a job being chosen.

We now move on to describe and estimate a job choice model to unwrap this evidence of jobs with female managers being chosen more on average but also incorporate the data on compensating differentials. We do this through the lens of worker preferences and the way beliefs operate in the absence of information on manager mentorship skill. The model also allows us to estimate our parameters as percentages of average annual wages to provide a better interpretation of their importance.

4 Model

Individuals are indexed by $i \in \{1, \dots, N\}$, and jobs are indexed by $j \in \{1, \dots, J\}$. Let X_j denote a K -dimensional vector of attributes of job j over which individuals have preferences. The utility of an individual i from job j is given by

$$U_{ij} = u_i(X_j) + \epsilon_{ij},$$

where ϵ_{ij} denotes all unobservables that affect the utility of individual i from job j . Individuals form expected utilities while reporting their job choice and the corresponding compensating differentials that would make them indifferent between jobs in expectation.

Individuals have preferences over working for a male manager (G), annual wages (W), availability of flexible hours (H) and manager mentorship rating (R). We denote this set of attributes as $X \equiv \{G, W, H, R\}$. In the complete scenarios, respondents observe X for each job. In the incomplete scenario, respondents observe \tilde{X} , where $\tilde{X} \equiv X \setminus R$. In the incomplete scenarios, when individuals do not observe the mentorship rating R , they use their beliefs on R given \tilde{X} to form their expected utilities.

The model is nonparametrically identified up to the distribution of $\epsilon_i \equiv \{\epsilon_{i1}, \dots, \epsilon_{iJ}\}$, as shown in Appendix A.1. In the following sections, to keep things simple, we use a linearly separable model. This functional form has no bearing on the identification of the parameter or its distribution, except for the additive separability of the error terms ϵ_{ij} .

The utility of an individual i with preference parameter vector $\beta_i \in \mathbb{R}^K$ from job j with K dimensions of attributes X_j is given by

$$U_{ij} = X_j' \beta_i + \epsilon_{ij} \quad (2)$$

Identification of more variants of the model allowing for various interactions is shown in Appendix A.3.

4.1 Complete scenarios

In the complete scenarios, individuals observe all attributes in set X for each job. The expected utility of individual i from job j conditional on its observable attributes in the complete scenarios is given by

$$\mathbb{E}_i[U_{ij}|X_j] = \sum_{x \in X} \beta_i^x x_j + \mathbb{E}_i(\epsilon_{ij}|X_j) \quad (3)$$

We assume that all individuals know their preferences and hence do not take expectations over them. As explained above, this draws a clear parallel with asking for choices instead of choice probabilities.

4.2 Incomplete scenarios

In the incomplete scenarios, the rating of the manager is mentioned but the data are shown as unavailable to the respondents. Hence, respondents form expectations over

them in reporting their choices and compensating differentials. Denote the set of observable attributes in job j as $\tilde{X}_j \equiv X_j \setminus \{R_j\}$ in the incomplete scenarios. The expected utility of individual i from job j conditional on its observable attributes \tilde{X}_j in the incomplete scenarios is given by

$$\mathbb{E}_i[U_{ij}|\tilde{X}_j] = \sum_{x \in \tilde{X}} \beta_i^x x_j + \mathbb{E}_i(\beta_i^R R_j + \epsilon_{ij}|\tilde{X}_j) \quad (4)$$

To proceed, we parameterize beliefs conditional on observables. Again, to keep things simple, we specify the process of beliefs linearly. We revisit this specification later and consider alternatives when we interpret how beliefs about manager ratings differ by manager gender.

$$\mathbb{E}_i(R_j|\tilde{X}_j) = \sum_{x \in \tilde{X}} \alpha_i^x x_j + \mathbb{E}_i(\eta_{ij}|\tilde{X}_j) \quad (5)$$

Observe that $\alpha_i^G = \mathbb{E}_i(R_j|G_j = \text{male}, W_j, H_j) - \mathbb{E}_i(R_j|G_j = \text{female}, W_j, H_j)$ represents how much on average individual i believes a male manager's mentorship rating differs from that of a female manager. Simplifying, the expected utilities in the incomplete scenarios are given by

$$\mathbb{E}_i[U_{ij}|\tilde{X}_j] = \sum_{x \in \tilde{X}} (\beta_i^x + \beta_i^R \alpha_i^x) x_j + \mathbb{E}_i(\epsilon_{ij} + \beta_i^R \eta_{ij}|\tilde{X}_j) \quad (6)$$

4.2.1 Discussion on the model: Choices and compensating differentials vs. choice probabilities

The hypothetical choice methodology literature has focused primarily on asking for respondents' choice probabilities if they were presented with a given choice at some point in the future. Hence, asking for probabilities makes sense when individuals are asked about future choices, and it is the resolvable uncertainty between today and the future that is the source of the probabilities. The difficulty is that the role of unobservables in this setting could be interpreted in one of two possible ways. Asking for probabilities implies resolution of resolvable uncertainty (Blass, Lach & Manski 2010). The uncertainty could arise from individuals learning about their preferences over time (Delavande & Manski 2015).²² However, this is at odds with how the model operates, whereby it is taken as given that individuals know their preferences and the unobservables are interpreted as preference shocks. Thus, it is not clear how individuals understand the instruction that

²²In Delavande & Manski (2015), voting behavior today could potentially be different from voting behavior in the future if the individual learns about her preferences over time prior to voting.

they should consider the attributes mentioned in the survey to be the only attributes along which the jobs vary.

Unlike in these models wherein ϵ_{ij} has an interpretation of resolvable uncertainty, in our model, ϵ_{ij} contains unobserved attributes that the worker may care about but does not observe. This derives from our asking respondents to make a job choice and report compensating differentials instead of choice probabilities. While we acknowledge that future choices maybe impacted by both resolvable (and unresolvable) uncertainties, to identify a model with both beliefs and preferences like ours without noninnocuous normalizations, we have to ask individuals to report their answers as if they are making the choices today. To reduce the impact of uncertainty resolutions, if any, we use a sample of individuals who are close to entering the labor market (one year away from graduating).

It is important to highlight that the role of ϵ_{ij} in the identification of our model is remarkably similar to its role in the models with resolvable uncertainties. In both cases, the assumption is that ϵ_{ij} is additively separable, and thus, the distribution of preferences (and in our case of beliefs as well) is nonparametrically identified up to the distribution of ϵ_{ij} . This is because whatever the econometrician does not observe—be it today’s unobservables or future uncertainties—needs to be held fixed through the instructions. Otherwise, there is nothing in the data that can help disentangle any effect of these unobservables if they interact with the observables. In this respect, we think that it is easier to instruct individuals to hold all else fixed instead of instructing them to hold the distribution of resolvable uncertainty fixed within scenarios. Future work can provide evidence on whether this distinction has any consequences for responses.

Hence, our instructions to individuals to hold all else fixed while reporting their choices and compensating differentials facilitate a clear link between the instructions provided and the role of the unobservables in the model. Given the instructions provided in our setting, additively separable unobservables do not vary across jobs within scenarios, conditional on observables. Additionally, asking for choices is more straightforward than asking for probabilities. Thus, it is more appealing if an effect of cognitive load on respondents’ survey responses is a concern.²³

5 Identification

In this section, we show how our experimental panel data on choices and compensating differentials identify the preference and belief parameters of our model of job choice by

²³We administered pilots where we asked for probabilities of choosing each job. Debriefs revealed that individuals spent longer understanding and/or calculating probabilities.

exploiting variation in the reported compensating differentials within and between the complete and incomplete scenarios.²⁴ As shown in Appendix Figure 3, individuals were instructed to assume that

Assumption (1): All attributes not mentioned in the survey were the same for all jobs.

Assumption (2): The reported compensating differential would increase only wages and change nothing else about the job.

Observe that instruction 1 is an assumption between jobs, while instruction 2 applies within jobs. The purpose of these instructions was to ensure that there was no selection on attributes not mentioned in the survey. Wiswall & Zafar (2018) delineate the importance of assumption (1) in a set-up such as our own in contrast to the settings in audit studies on discrimination, where there is little preventing employers from making assumptions about job applicants conditional on the observables in their resumes.²⁵ For both scenarios, assumption (2) implies that the compensating differential increases only the wage and does not change the conditional expectation of the unobservables. Note that for the incomplete scenarios, it applies to the conditional expectation of the manager rating as well, as we show in the following sections. We use the data on compensating differentials to equate the expected utilities in the complete and incomplete scenarios. These two assumptions, which form a clear parallel to the instructions given to the respondents, form the basis of our identification strategy.

5.1 Identification of preferences from the complete scenarios

In this section, we show the identification of the preference parameters β_i^x for all $x \in \{G, H, R\}$. The parameter of interest is β_i^G , which is the preference for male managers over female managers. We use assumptions (1) and (2) to identify the preferences using the variation in compensating differentials within the complete scenarios.

Implication of Assumption (1):

The instructions imply that unobservables across different jobs in conditional expectations are the same within each scenario. For every individual i and every job $j \neq k$ within each complete scenario,

²⁴In the Appendix, we also write a more flexible model where the rating variable is used as a signal for overall manager quality and show the identification in that setting.

²⁵For more on the use of audit studies in detecting discrimination, see Heckman (1998) and related papers within it.

$$\mathbb{E}_i(\epsilon_{ij}|X_j) = \mathbb{E}_i(\epsilon_{ik}|X_k)$$

Implication of Assumption (2):

In the complete scenarios, for each job k , individuals observe the vector of attributes $X_k = \{G_k, H_k, W_k, R_k\}$. Suppose that individual i chooses job j and then provides a compensating differential of Δ_{ijk} that she would require to choose job k instead. The instructions imply that for unchosen jobs such as job k ,

$$\mathbb{E}_i(\epsilon_{ik}|X_k, \Delta_{ijk}) = \mathbb{E}_i(\epsilon_{ik}|X_k) \quad (7)$$

Given the above, equating the expected utilities between job j and job k with the provided compensating differential of Δ_{ijk} and normalizing $\beta_i^W = 1$, we have

$$\begin{aligned} \mathbb{E}_i(U_{ij}|X_j) &= \mathbb{E}_i(U_{ik}|X_k, \Delta_{ijk}) \\ \Delta_{ijk} &= \sum_{x \in X} \beta_i^x (x_j - x_k) \end{aligned} \quad (8)$$

Thus, the preference parameters β_i^x are identified using the variation in the reported compensating differentials in the complete scenarios under assumptions (1) and (2) for all $x \in X$.

5.2 Identification of beliefs from the complete and incomplete scenarios

Now, we turn to the incomplete scenarios in conjunction with the complete scenarios and show the identification of the belief parameter α , exploiting the variation between the reported compensating differentials between the complete and incomplete scenarios.

Implication of assumption (1):

The instruction implies that unobservables affecting utilities and beliefs across different jobs in conditional expectations are the same within each scenario. For every individual i and every job $j \neq k$ within each incomplete scenario,

$$\begin{aligned} \mathbb{E}_i(\epsilon_{ij}|\tilde{X}_j) &= \mathbb{E}_i(\epsilon_{ik}|\tilde{X}_k) \\ \mathbb{E}_i(\eta_{ij}|\tilde{X}_j) &= \mathbb{E}_i(\eta_{ik}|\tilde{X}_k) \end{aligned}$$

Implication of Assumption (2):

In the incomplete scenarios, for each job k , individuals observe the vector of attributes $\tilde{X}_k = \{G_k, H_k, W_k\}$. Suppose that individual i chooses job j and provides a compensating differential of $\tilde{\Delta}_{ijk}$ that she would require to choose job k instead. All the compensating

differential does is increase the wages in job k by $\tilde{\Delta}_{ijk}$. The implication of assumption (2) is that it has no effect on the conditional expectation of managers' ratings or on the conditional expectation of the unobservables affecting utility. That is,

$$\begin{aligned} \mathbb{E}_i(R_k|\tilde{X}_k, \tilde{\Delta}_{ijk}) &= \mathbb{E}_i(R_k|\tilde{X}_k) \\ \text{and } \mathbb{E}_i(\epsilon_{ik}|\tilde{X}_k, \tilde{\Delta}_{ijk}) &= \mathbb{E}_i(\epsilon_{ik}|\tilde{X}_k) \end{aligned} \quad (9)$$

Thus, the expected utility from job k taking into account the compensating differential of $\tilde{\Delta}_{ijk}$ along with the normalization of β_i^W to 1 is

$$\begin{aligned} \mathbb{E}_i(U_{ik}|\tilde{X}_k, \tilde{\Delta}_{ijk}) &= \beta_i^G G_k + \beta_i^H H_k + (W_k + \tilde{\Delta}_{ijk}) + \beta_i^R \mathbb{E}_i(R_k|\tilde{X}_k, \tilde{\Delta}_{ijk}) + \mathbb{E}_i(\epsilon_{ik}|\tilde{X}_k, \tilde{\Delta}_{ijk}) \\ &= \beta_i^G G_k + \beta_i^H H_k + (W_k + \tilde{\Delta}_{ijk}) + \beta_i^R \mathbb{E}_i(R_k|\tilde{X}_k) + \mathbb{E}_i(\epsilon_{ik}|\tilde{X}_k) \\ &= \sum_{x \in \tilde{X}} (\beta_i^x + \beta_i^R \alpha_i^x) x_k + \tilde{\Delta}_{ijk} + \mathbb{E}_i(\epsilon_{ik} + \beta_i^R \eta_{ik}|\tilde{X}_k) \end{aligned} \quad (10)$$

given that beliefs about ratings are modeled as in equation (5).

Individuals are assumed to know their preference parameters and hence not to take expectations over them. Equating the expected utilities between job j and job k with the provided compensating differential of $\tilde{\Delta}_{ijk}$ and as before normalizing $\beta_i^W = 1$ under A1, we have

$$\begin{aligned} \mathbb{E}_i(U_{ij}|\tilde{X}_j) &= \mathbb{E}_i(U_{ik}|\tilde{X}_k, \tilde{\Delta}_{ijk}) \\ \tilde{\Delta}_{ijk} &= \sum_{x \in \tilde{X}} \underbrace{(\beta_i^x + \beta_i^R \alpha_i^x)}_{\tilde{\beta}_i^x} (x_j - x_k) \end{aligned} \quad (11)$$

Now, identification of α is straightforward. Denote for each attribute $x \in \tilde{X}$

$$\tilde{\beta}_i^x \equiv \beta_i^x + \beta_i^R \alpha_i^x \quad (12)$$

Observe that $\tilde{\beta}^x$ is comprised of two terms: the preference parameter for attribute x and how much x affects the belief about the manager's rating, given that the individual cares about the manager's rating.

Thus, given the identification of β_i^x and $\tilde{\beta}_i^x$, we have for all $x \in \tilde{X}$

$$\alpha_i^x = \frac{\tilde{\beta}_i^x - \beta_i^x}{\beta_i^R} \quad (13)$$

Thus, α_i^x is identified $\forall x \in \tilde{X}$ as long as $\beta_i^R \neq 0$

We want to end this section with a small discussion on the intuitive difference between the two sets of equations on compensating differentials in the complete and incomplete scenarios.

In the complete scenarios, the compensating differentials are a function of how jobs j and k vary in their attributes, weighted by how much individual i cares about each of those attributes. In contrast, in the incomplete scenarios, they are a function of how jobs j and k vary in their attributes apart from mentorship, weighted by not only how much individuals care about each attribute but also by how much they believe each attribute is correlated with mentorship skill—very much in the spirit of omitted variable bias.

Additionally note that there are two circumstances when beliefs are not identified. Observe that α_i^x is not identified if $\beta_i^R = 0$, i.e., when i does not care about mentorship ability. The underlying intuition is that if individuals do not care about manager mentorship, then any belief distribution can rationalize the observed data. This is because the variation in the observed choices and compensating differentials are independent of mentorship skill. Secondly, beliefs are not identified if individuals believe mentorship ability is independent of all observed attributes. In particular, in that case, we would have for each individual i , $\mathbb{E}_i(R_k|\tilde{X}_k) = \mathbb{E}_i(R_k)$ for all jobs k , which is a constant, though it could vary by i . However, since it does not vary with the observed job attributes, any within-individual variation cannot be used to identify beliefs.

6 Estimation

We use variation in the reported compensating differentials to estimate the preference and belief parameters of our model. The compensating differentials are reported with two different but independent measurement errors. One measurement error results from the reporting of the compensating differentials in multiples of five, as shown in Figure 4. This is also observed in surveys asking for choice probabilities (Blass, Lach & Manski 2010). The second type of measurement error arises from the design of the survey. Individuals had a slider to report their compensating differentials, and sliding the slider could cause some random measurement error, even though individuals could see the exact amount as they slid the slider. If both types of measurement errors are classical in nature, they will only inflate the standard errors. The consistency of the estimates is not affected.

Denote as Δ_{ijk}^* and $\tilde{\Delta}_{ijk}^*$ the latent compensating differentials and as e_i and \tilde{e}_i the composite classical measurement errors in the complete and incomplete scenarios, respectively. Thus, we have the following set of estimating equations for each individual i and

each pair of jobs j and k within every scenario:

$$\begin{aligned}\Delta_{ijk}^* &= \sum_{x \in X} \beta_i^x (x_j - x_k) + e_i \\ \tilde{\Delta}_{ijk}^* &= \sum_{x \in \tilde{X}} (\beta_i^x + \beta_i^R \alpha_i^x) (x_j - x_k) + \tilde{e}_i\end{aligned}\tag{14}$$

With classical measurement errors, we have

$$\begin{aligned}\mathbb{E}[\Delta_{ijk}^* | X_j, X_k] &= \sum_{x \in X} \beta_i^x (x_j - x_k) = (X_j - X_k)' \beta_i \\ \mathbb{E}[\tilde{\Delta}_{ijk}^* | \tilde{X}_j, \tilde{X}_k] &= \sum_{x \in \tilde{X}} (\beta_i^x + \beta_i^R \alpha_i^x) (x_j - x_k) = (\tilde{X}_j - \tilde{X}_k)' \tilde{\beta}_i\end{aligned}\tag{15}$$

We jointly estimate the above system of equations (14) using constrained least squares where we normalize the preference parameter on wages to one ($\beta_i^W = 1$). By construction, the estimating equations have no constant since they are derived by equating utility functions. We use the block bootstrap at the respondent level to allow for arbitrary correlation among responses within each respondent. We describe the details of the joint estimation and the bootstrap algorithm in Appendix A.2.

6.1 Model estimates

In this section, we discuss the estimates of the preference and belief parameters of the job choice model described in the previous section. Estimates of alternate models allowing for nonlinearities, presented in Appendix A.3, have similar results and are presented in the Online Appendix. In the first subsection, we discuss the estimates for preference parameters focusing on the preference on manager gender. Next, we discuss the preference and thus demand for good mentors. We discuss estimates of the belief parameters focusing on beliefs on mentorship ability by manager gender in the final subsection. This order of discussion is deliberate because only after establishing how much individuals care about managers' gender and mentorship skill is it useful to discuss beliefs on mentorship by manager gender.

Since the utility parameter on wages is normalized to 1 and wages are in hundred thousand (lakh) INR, the estimates should be interpreted as valuations of each attribute in hundred thousand INR.²⁶ However, for even better interpretability, we also present these estimates by converting them into percentages of average annual wages.

²⁶Alongside the estimates, we present the purchasing power parity equivalent in USD.

6.1.1 Preference for working for female managers

Table 11 shows the estimates from equation system (14) representing indifferences in the complete and incomplete scenarios jointly exploiting the variation in reported compensating differentials across jobs with exogenously varying attributes. We first discuss the complete scenarios, from which we derive the pure preference parameter estimates on manager gender, flexibility of hours and manager mentorship rating, with the preference parameter on average wages normalized to 1.

The most striking result is the evidence of a strong preference to work for female managers as shown in the panel of complete scenarios. The preference parameter for male managers (β^G) is negative and statistically significant. On average, individuals are willing to give up 12 thousand INR (\approx \$570) to work for a female manager. This value corresponds to 1.7% of average annual wages. The 95% confidence interval²⁷ of this value as a percentage of average annual wages is from 1.3% to 2.2%.

For the incomplete scenarios, the “biased” estimate of the preference for male managers ($\tilde{\beta}^G$) is statistically indistinguishable from zero. The difference in the estimates across the complete (β^G) and incomplete scenarios ($\tilde{\beta}^G$) provides evidence that in the absence of information on manager mentorship quality, individuals believe that male managers are better mentors. The results show that beliefs and preferences operate in opposite directions to generate these indifferences. This difference in the estimates between the complete and incomplete scenarios is evidence of statistical discrimination against female managers if and only if individuals prefer to work for managers who are better mentors. We now turn to discussing preferences on manager mentorship ability.

6.1.2 Demand for mentorship quality in managers

As explained earlier in the identification section, if individuals do not care about manager mentorship (i.e., $\beta^R = 0$), then beliefs are fundamentally unidentified, and thus, this difference cannot be taken as evidence of statistical discrimination. Table 11 shows that respondents care about high-quality mentorship from their managers. Individuals on average are willing to give up approximately 11% of their average annual wages (\approx \$ 3,800 or 80 thousand INR) to work for a mentor who ranks one point higher in mentorship ability on a five-point scale. Converting this to standard-deviation units, given the variation in mentorship ratings in the scenarios presented to the respondent, a one-standard-deviation increase in mentorship skill is valued at 5.65% of average annual wages. This estimate could seem large at first glance. However, it is not surprising given that the re-

²⁷Bias-corrected bootstrap confidence interval.

spondents are jobseekers who are about to enter the labor market for the first time. Under diminishing returns to mentorship, a marginal increase in mentorship ability is of much higher value to first-time jobseekers than to experienced workers in the labor market.

A potential concern could revolve around how respondents interpret average mentorship quality if the respondent is very different from the current workers on whose ratings the mentorship quality data are based. In such a case, the mentorship variable could be uninformative to the respondent.^{28, 29} Even though our scenarios in the survey do not go into the specifics of types of jobs or industry to reduce cognitive load, if respondents were indeed thinking of jobs dominated by out-group workers such that the mentorship quality rating was uninformative, then we would observe evidence of this in the data. By contrast, the lower limit of the 95% confidence interval on the preference for high-rated mentors is at 10.5% of average annual wages, far from zero.³⁰

6.1.3 Beliefs

The estimates of the belief parameter on male managers' mentorship (α^G) are obtained from the estimates of the vectors $(\tilde{\beta}^G, \beta^G, \beta^R)$ using equation (13). The first row of Table 12 shows that on average male managers are believed to have a 0.14 points (on a 5-point scale) or 0.28 standard deviations higher mentorship rating than female managers. Classical measurement error in the reported compensating differentials leading to inflation of standard errors in the estimation of the preference parameters will trickle down to the standard errors of the belief parameters. Despite this, the belief parameter estimate is statistically significant. The estimates in conjunction with the model and the evidence that mentorship skill is a highly sought-after manager attribute imply that in the absence of a manager rating, both genders believe that male managers are better mentors than female managers.

The estimates of α^G are hard to interpret since they represent relative beliefs on a five-point scale. A more interpretable measure of beliefs—in monetary terms—is the valuation of beliefs on male manager mentorship ($\beta^R \alpha^G$). In Table 13, we report these estimates of the valuation of worker beliefs. Observe that the difference in the parameters between the complete and incomplete scenarios provides us exactly that: $\tilde{\beta}^x - \beta^x = \beta^R \alpha^x$ for all

²⁸Recall that the mentorship rating was presented in the survey as the average mentorship rating of the manager provided by his or her current workers.

²⁹An example could be a woman evaluating a job offer from the construction sector, which is heavily male-skewed.

³⁰There is less concern about other interpretations of mentorship. For example, it is unlikely that English majors would answer our questions by invoking upon themselves the extra cognitive load to think about jobs outside their domain of specialization.

$x \in \{G, W, H\}$. Thus, the estimates of $\beta^R \alpha^G$ provide us an estimate of the valuation of statistical discrimination.³¹

7 Heterogeneity

7.1 Distributions

Each respondent in our survey answered questions in 20 scenarios with 3 jobs each. This gives us 40 unique data points of compensating differentials across jobs of varying attributes for each respondent—20 from the incomplete and 20 from the complete scenarios. We use these data to estimate the model for each individual separately. Using the empirical distribution of respondent-specific estimates, we can compute the sample mean preferences, and given the standard deviation, we can compute the standard errors of the mean. More interestingly, however, using the empirical distributions, we carry out two informative exercises. First, we quantify the proportion of individuals who statistically discriminate against female managers. Second, we regress the estimated parameters on sample characteristics to obtain the correlations between discrimination and observed characteristics.^{32,33} We also address the concern that the individual parameters may be estimated with noise and discuss bounds on the estimates under reasonable assumptions. For better interpretability of the figures, we have converted the preference and belief parameters into percentages of average annual wages.

As shown in Figure 5, the empirical CDF reveals that in the absence of information on mentorship quality, 60% of respondents believe that female managers are worse mentors than male managers. If the true underlying distributions of preferences, beliefs and noise are symmetric, then this number is a lower bound on the proportion of individuals who statistically discriminate against female managers. This is because under symmetry, the median is unaffected, and hence, the true cumulative distribution will intersect zero at a lower point than what we see in Figure 5. Upon comparison of the estimate of the mean

³¹In addition, note that $\frac{\mathbb{E}(\tilde{\beta}_i^G - \beta_i^G)}{\mathbb{E}(\beta_i^R)} \neq \mathbb{E}\left(\frac{\tilde{\beta}_i^G - \beta_i^G}{\beta_i^R}\right) = \mathbb{E}(\alpha_i^G)$. However, $\mathbb{E}(\tilde{\beta}_i^G - \beta_i^G) = \mathbb{E}(\beta_i^R \alpha_i^G)$ provides the average valuations of beliefs.

³²The parameter estimates that we regress on observed demographics could be noisy, being estimated from 40 observations. This will overstate the standard errors of the regression coefficients. A finding of statistically significant estimates despite the overstated standard errors would make an even stronger case for heterogeneity.

³³One could also take the route of testing for heterogeneity using wild bootstrap- t (Cameron, Gelbach & Miller (2008), Busso, Gregory & Kline (2013)). However, in our case, the test is nonstandard since it is testing at the boundary of the parameter space—i.e., testing the null hypothesis of zero variance against the alternative hypothesis of positive variance.

and its standard error obtained by estimating the model for each respondent with the base model with all respondents together, we find that their values are very close. This would not be the case if the noise on average were not zero, even though the individual estimates could still be noisy. Similarly, in Figure 6, we observe that upon receiving information on manager quality, at least 62% of individuals prefer to work for female managers. In addition, it is important to note that it is not everyone who prefers to work for a female manager, even though on average we find a preference to work for female managers. The distribution reveals two important points. First, there are more individuals who prefer to work for female managers. Second, the amount of money that would make an individual who prefers to work for a female manager switch to a job with a male manager is higher than the amount needed to make an individual who prefers to work for a male manager switch to a job with a female manager.

7.2 Correlates of parameters with observable characteristics

In this section, we present results on how respondents' preference and belief parameters correlate with their observable characteristics. We regress each individual's estimated parameters (preference to work for male managers (β_i^G) and valuation of beliefs on male managers' mentorship skill relative to female managers' ($\beta_i^R \alpha_i^G$)) on the respondent's gender, major, family income, parental education, and parental employment. We estimate the following equations:

$$\begin{aligned} \hat{\beta}_i^G = & \gamma_0 + \gamma_1 \text{Gender}_i + \gamma_2 \text{FamilyIncome}_i + \gamma_3 \text{Major}_i + \\ & \gamma_5 \mathbb{I}(\text{Mother more educated than Father})_i + \gamma_6 \text{EmployedFather}_i + \gamma_7 \text{EmployedMother}_i + \\ & \gamma_8 \mathbb{I}(\text{Mother more educated than Father})_i * \text{EmployedFather}_i * \text{EmployedMother}_i + u_i \end{aligned}$$

and

$$\begin{aligned} \widehat{\beta_i^R \alpha_i^G} = & \phi_0 + \phi_1 \text{Gender}_i + \phi_2 \text{FamilyIncome}_i + \phi_3 \text{Major}_i + \\ & \phi_5 \mathbb{I}(\text{Mother more educated than Father})_i + \phi_6 \text{EmployedFather}_i + \phi_7 \text{EmployedMother}_i + \\ & \phi_8 \mathbb{I}(\text{Mother more educated than Father})_i * \text{EmployedFather}_i * \text{EmployedMother}_i + e_i \end{aligned}$$

We remove respondents whose estimated valuation of mentorship skill is either close to zero or negative. Our rationale for removing respondents with zero (or close to zero)

valuations of mentorship skill comes from the identification argument. Any belief distribution can be used to rationalize responses of respondents who do not value mentorship skill, and thus, beliefs are fundamentally unidentified for such respondents. We also drop respondents whose valuation of manager mentorship ability is negative. This is because it is not clear how one can interpret a negative valuation of mentorship ability. This leads us to drop 43 such respondents and brings our sample size to 535. We present these results in Table 14. The first column has the estimates of the correlation of respondents' characteristics with their preferences for male managers and the second column the estimates of the correlation of the same set of characteristics with their valuation of beliefs on male managers' mentorship ability relative to female managers'.

We find evidence of the underlying heterogeneity in preferences and beliefs being correlated with individual education, maternal education, and parental employment. We find that respondents with a major in engineering are more likely to prefer to work for female managers than those with major in the humanities. Respondents whose mothers are weakly more educated than their fathers are less likely to have negative beliefs about female managers. Respondents with employed fathers are more likely to have negative beliefs on female manager mentorship skill. Relative to respondents whose mothers are less educated than their fathers and respondents whose both parents are unemployed, respondents with at least employed fathers (irrespective of the mother's employment status) are less likely to have negative beliefs about female manager mentorship ability.

The correlations of the characteristics above with the underlying heterogeneity in the distribution of beliefs on female managers point toward an important distinction between the origins of gender and race discrimination based on beliefs. Individuals who have lived in race-segregated neighborhoods have very little chance to learn about people of other races not represented in the neighborhood. However, growing up without a mother is not as common as living in a race-segregated neighborhood.³⁴ This argument finds support in Alesina et al. (2021), who find evidence that individuals who have lived in racially diverse neighborhoods tend not to hold biased beliefs on individuals of other races.

³⁴We thank Martha Bailey for sharing this observation.

8 Robustness

8.1 Validity of belief parameter estimates

In the penultimate section of the survey, we directly elicited beliefs on manager mentorship. We did this to compare the results with those obtained from our information experiment. The objective of this exercise was not to see how close the two estimates of α^G are. As discussed earlier, the estimates of α from the model could be noisy for individuals whose β^R is close to zero. Moreover, we cannot rule out the possibility that learning could impact the estimates on beliefs obtained from the direct belief elicitation, which could thus turn out to be different from the estimates obtained from the information experiment within the job choice scenarios. The objective was instead to see whether the result on the belief in the inferiority of female managers as mentors held when beliefs were directly elicited. Although we recognize that direct elicitation of beliefs on mentorship by manager gender might induce social desirability bias, it was nevertheless a useful exercise to see to what extent the results might diverge from those obtained from our information experiment.

This part of the survey presented respondents with 10 jobs. Each job description exogenously varied the manager’s first name, wages and flexibility of hours. Alongside each job, we provided a zero-to-five slider scale and asked respondents to report their expected rating for each manager in each of the jobs. On these data, we project a linear model of reported expected ratings on manager gender, annual wages and flexibility of hours in alignment with the parametrization of beliefs in the job choice model.

Table 15 reports the estimates for all individuals and by respondent gender. We see that this exercise leads to the same qualitative conclusion on the belief parameter estimated from the information experiment data. We do observe that men’s directly elicited beliefs are substantially lower than those elicited from the information experiment. This could arise from two sources that we cannot distinguish—shading of directly reported beliefs due to social desirability bias or learning from the scenarios that male and female managers have similar mentorship ratings. We do not observe this stark difference among female respondents. Note that these data alone can help us estimate only α^x and not $\beta^r \alpha^x$ for all $x \in \{G, H, W\}$. However, the primary objective of these data is served in that they corroborate that the sign of α^G is positive in the estimates obtained both from the job choice model using the choice and compensating differential data and from the directly elicited belief data. This is an important result because it gives us additional evidence that individuals believe that women have lower average mentorship quality. Hence, any concerns about individuals manipulating their responses and misreporting their preference

to work for female managers is not an issue given this corroborating evidence.

8.2 Further checks

A concern is whether individuals followed the instructions provided to them at the beginning of the survey. To evaluate this, we designed specific questions, both direct and indirect, at the end of the survey to infer whether individuals followed the instructions. We dropped the 2.2% of the sample respondents who failed these checks and reestimated the model. The estimates are robust to the use of this restricted sample.

To deal with survey inattention, we also dropped the 1% of sample respondents who finished the survey in less than 15 minutes (Mas & Pallais 2017). The choice of 15 minutes was motivated by the distribution of time completion of the survey—the 1st percentile of completion duration was at 13.89 minutes. Our estimates are robust to the imposition of this restriction as well.

9 Additional Discussion and Implications

Given our results in the incomplete scenarios, it is important to note that an observation of indifference between male and female managers in the data should not necessarily be interpreted as evidence of no preference for one gender over the other. This can happen when preferences and beliefs operate in opposite directions, as they do in our incomplete scenarios. We cannot comment on whether such beliefs held by individuals are biased. To do so, we would need data on the population distributions of mentorship skill by manager gender. Although we are limited in this respect, we would like to highlight literature across various contexts showing consistent evidence that individuals in general do not have a good sense of population distributions (e.g., Wiswall & Zafar (2015), Bordalo, Coffman, Gennaioli & Shleifer (2016), Alesina, Miano & Stantcheva (2019), Bursztyn et al. (2018), Bohren, Imas & Rosenberg (2019), Alesina & Stantcheva (2020), Hvidberg, Kreiner & Stantcheva (2020), Bleemer & Zafar (2018)).

There are two primary implications of our research. First, if we fail to consider labor supply-side selection, we cannot obtain a complete picture of group-level inequalities. Additionally studying beliefs and preferences on the supply side is essential because these affect search and equilibrium matches. As discussed in the introduction, this is particularly important in tight labor markets or markets where workers are in general harder to replace. The second implication is how firms might respond differently to such worker beliefs and preferences, which could generate different rates of promotion of women to

managerial positions.³⁵

These implications provide multiple potential avenues for future research extending from our paper. If firms have strong priors that matching workers with their preferred managers increases match productivity, it could potentially lead to women being promoted at higher rates, conditional on mentoring capabilities. However, women could still be promoted at lower rates if firms have sufficiently discriminatory preferences against women and are willing to forgo the increased profits resulting from more efficient matches. Thus, in a way, worker preferences could be used to test for discriminatory practices by the firm executives who decide whom to promote to managerial positions. When workers prefer to work for female managers and, conditional on productivity, women are still promoted at lower rates than men, it could be interpreted as evidence of discrimination. It would be also interesting to explore whether individuals' preferences and perceptions about average preferences differ and, more importantly, whether their perceptions are incorrect. In the spirit of [Bursztyn, González & Yanagizawa-Drott \(2018\)](#), incorrect perceptions of preferences could be an additional reason why female managers are promoted at lower rates.³⁶ Another avenue for future research is to explore whether having more female managers allows firms to profitably compete for workers with otherwise-similar firms if there is overall preference to work for female managers. Additionally, if there are information asymmetries between incumbent and competing firms ([Pinkston \(2009\)](#), [Kahn \(2013\)](#)), then there are even higher information rents to be taken advantage of.

There are many other instances in which our research applies. Examples include job-seekers who have alumni networks in firms and can obtain information on managers and manager quality, thus affecting their search and consequently final match. Another example is individuals who are already employed but are seeking to switch teams within firms.

10 Conclusion

In this paper, we provide novel documentation of discrimination by workers on manager's gender in choosing jobs. To do so, we designed and conducted a hypothetical job

³⁵[Cai & Wang \(2022\)](#) show that firms and supervisors do respond to worker feedback in a firm-wide field experiment.

³⁶[Bursztyn, González & Yanagizawa-Drott \(2018\)](#) show that in Saudi Arabia, husbands individually prefer having their wives participate in the labor force but misperceive social norms and believe that such preferences are uncommon on average. When their misperceptions are corrected, husbands enroll their wives in a costly training program, thus increasing female labor force participation.

choice survey involving a novel information experiment among jobseeking students at a highly selective university in India. We presented respondents with a series of hypothetical job scenarios consisting of jobs with exogenously varying attributes (annual wages, flexibility of hours, and manager name and mentorship rating). Respondents were asked to choose their most preferred job and report the wage compensating differentials that would make them indifferent between jobs. In the first ten scenarios (incomplete scenarios), individuals were shown jobs with all attributes except the manager mentorship rating. In the last ten scenarios (complete scenarios), individuals observed all of the above attributes for each job. We identify preferences using the variation in compensating differentials within the complete scenarios. The variation in compensating differentials between the complete and incomplete scenarios provides us with the necessary variation to identify beliefs on mentorship skill. We find that in the absence of information on manager mentorship, where their choices are driven by both preferences and beliefs, workers are indifferent between male and female managers. However, in the presence of information on manager mentorship skill, we find a strong preference to work for female managers. Using a structural model of job choice, we estimate that individuals are willing to forgo on average 1.7% of average annual wages to work for female managers. Hence, in the absence of additional information on manager mentorship ability, female managers are believed to be worse mentors than male managers. We quantify these negative beliefs against female managers at approximately 1.6% of average annual wages. We corroborate these results on negative beliefs on female manager mentorship using additional data on directly elicited beliefs. Importantly, we uncover rich heterogeneity in the underlying distributions of individuals' preferences and beliefs. Estimating the model for each worker, we find that preferences and beliefs correlate with demographics in the expected directions. In particular, we do not find evidence of negative beliefs on female manager mentorship ability among female respondents majoring in engineering or among those whose mothers are more educated than their fathers. As discussed in detail in the previous section, our results suggest that estimates of discrimination by firm executives in generating a glass ceiling for women at managerial levels could be downward biased if, conditional on quality, women are still promoted at lower rates. To the extent that this is the case, our paper sheds light on how additional data on worker preferences and beliefs on manager characteristics could be used to indirectly test for discrimination among the firm executives who decide on promotions.

References

- Abel, M. (2019), Do workers discriminate against female bosses? Institute of Labor Economics (IZA) Discussion Paper No. DP12611.
- Abel, M. & Buchman, D. (2020), The effect of manager gender and performance feedback: Experimental evidence from india. Institute of Labor Economics (IZA) DP No. 13871.
- Acemoglu, D., He, A. & le Maire, D. (2022), Eclipse of rent-sharing: The effects of managers' business education on wages and the labor share in the us and denmark, Technical report, National Bureau of Economic Research.
- Adams-Prassl, A. & Andrew, A. (2019), Preferences and beliefs in the marriage market for young brides. CEPR Discussion Paper No. DP13567.
- Agan, A. & Starr, S. (2018), 'Ban the box, criminal records, and racial discrimination: A field experiment', *The Quarterly Journal of Economics* **133**(1), 191–235.
- Alam, M. M. U. (2019), Identification in models of discrimination. Working Paper.
- Alesina, A., Ferroni, M. F. & Stantcheva, S. (2021), Perceptions of racial gaps, their causes, and ways to reduce them, Working Paper 29245, National Bureau of Economic Research.
URL: <http://www.nber.org/papers/w29245>
- Alesina, A., Miano, A. & Stantcheva, S. (2019), Immigration and redistribution. National Bureau of Economic Research Working Paper 24733.
- Alesina, A. & Stantcheva, S. (2020), 'Diversity, immigration, and redistribution', *AEA Papers and Proceedings* **110**, 329–34.
- Altonji, J. G. & Pierret, C. R. (2001), 'Employer learning and statistical discrimination', *The Quarterly Journal of Economics* **116**(1), 313–350.
- Ameriks, J., Briggs, J., Caplin, A., Shapiro, M. D. & Tonetti, C. (2020), 'Long-term-care utility and late-in-life saving', *Journal of Political Economy* **128**(6), 2375–2451.
- Arrow, K. J. (1972), Some mathematical models of race discrimination in the labor market, in A. Pascal, ed., 'Racial discrimination in Economic Life', Lexington, Mass.: Lexington Books, pp. 187–204.

- Arrow, K. J. (1998), 'What has economics to say about racial discrimination?', *Journal of Economic Perspectives* **12**(2), 91–100.
- Athey, S., Avery, C. & Zemsky, P. (2000), 'Mentoring and diversity', *American Economic Review* **90**(4), 765–786.
- Ayalew, S., Manian, S. & Sheth, K. (2021), 'Discrimination from below: Experimental evidence from ethiopia', *Journal of Development Economics* **151**, 102653.
- Becker, G. S. (1971), *The economics of discrimination*, Second Edition, Chicago: University of Chicago Press.
- Benson, A., Li, D. & Shue, K. (2019), 'Promotions and the peter principle', *The Quarterly Journal of Economics* **134**(4), 2085–2134.
- Bertrand, M. & Duflo, E. (2017), Field experiments on discrimination, in E. D. Abhijit Vinayak Banerjee, ed., 'Handbook of economic field experiments', Vol. 1, Elsevier, pp. 309–393.
- Bertrand, M. & Mullainathan, S. (2004), 'Are emily and greg more employable than lakisha and jamal? a field experiment on labor market discrimination', *American Economic Review* **94**(4), 991–1013.
- Blass, A. A., Lach, S. & Manski, C. F. (2010), 'Using elicited choice probabilities to estimate random utility models: Preferences for electricity reliability', *International Economic Review* **51**(2), 421–440.
- Blau, F. D., Currie, J. M., Croson, R. T. & Ginther, D. K. (2010), 'Can mentoring help female assistant professors? interim results from a randomized trial', *American Economic Review* **100**(2), 348–52.
- Blau, F. D. & Kahn, L. M. (2017), 'The gender wage gap: Extent, trends, and explanations', *Journal of economic literature* **55**(3), 789–865.
- Bleemer, Z. & Zafar, B. (2018), 'Intended college attendance: Evidence from an experiment on college returns and costs', *Journal of Public Economics* **157**, 184–211.
- Bohren, J. A., Imas, A. & Rosenberg, M. (2019), 'The dynamics of discrimination: Theory and evidence', *American Economic Review* **109**(10), 3395–3436.

- Boneva, T., Buser, T., Falk, A., Kosse, F. et al. (2021), The origins of gender differences in competitiveness and earnings expectations: Causal evidence from a mentoring intervention, Technical report. Working Paper No. 2021-049 , Human Capital and Economic Opportunity Working Group.
- Boneva, T. & Rauh, C. (2018), 'Parental beliefs about returns to educational investments - the later the better?', *Journal of the European Economic Association* **16**(6), 1669–1711.
- Bordalo, P., Coffman, K., Gennaioli, N. & Shleifer, A. (2016), 'Stereotypes', *The Quarterly Journal of Economics* **131**(4), 1753–1794.
- Boyer, M., De Donder, P., Fluet, C., Leroux, M.-L. & Michaud, P.-C. (2017), Long-term care insurance: Knowledge barriers, risk perception and adverse selection, Technical report, National Bureau of Economic Research.
- Bursztyn, L., González, A. L. & Yanagizawa-Drott, D. (2018), Misperceived social norms: Female labor force participation in saudi arabia. National Bureau of Economic Research Working Paper 24736.
- Busso, M., Gregory, J. & Kline, P. (2013), 'Assessing the incidence and efficiency of a prominent place based policy', *American Economic Review* **103**(2), 897–947.
- Cai, J. & Wang, S.-Y. (2022), 'Improving management through worker evaluations: Evidence from auto manufacturing', *The Quarterly Journal of Economics* **137**(4), 2459–2497.
- Cameron, A. C., Gelbach, J. B. & Miller, D. L. (2008), 'Bootstrap-based improvements for inference with clustered errors', *The Review of Economics and Statistics* **90**(3), 414–427.
- Charles, K. K. & Guryan, J. (2008), 'Prejudice and wages: an empirical assessment of beckers the economics of discrimination', *Journal of political economy* **116**(5), 773–809.
- Conlon, J. J., Pilossoph, L., Wiswall, M. & Zafar, B. (2018), Labor market search with imperfect information and learning, Technical report, National Bureau of Economic Research.
- Delavande, A. & Manski, C. F. (2015), 'Using elicited choice probabilities in hypothetical elections to study decisions to vote', *Electoral Studies* **38**, 28–37.
- Dey, M. S. & Flinn, C. J. (2005), 'An equilibrium model of health insurance provision and wage determination', *Econometrica* **73**(2), 571–627.

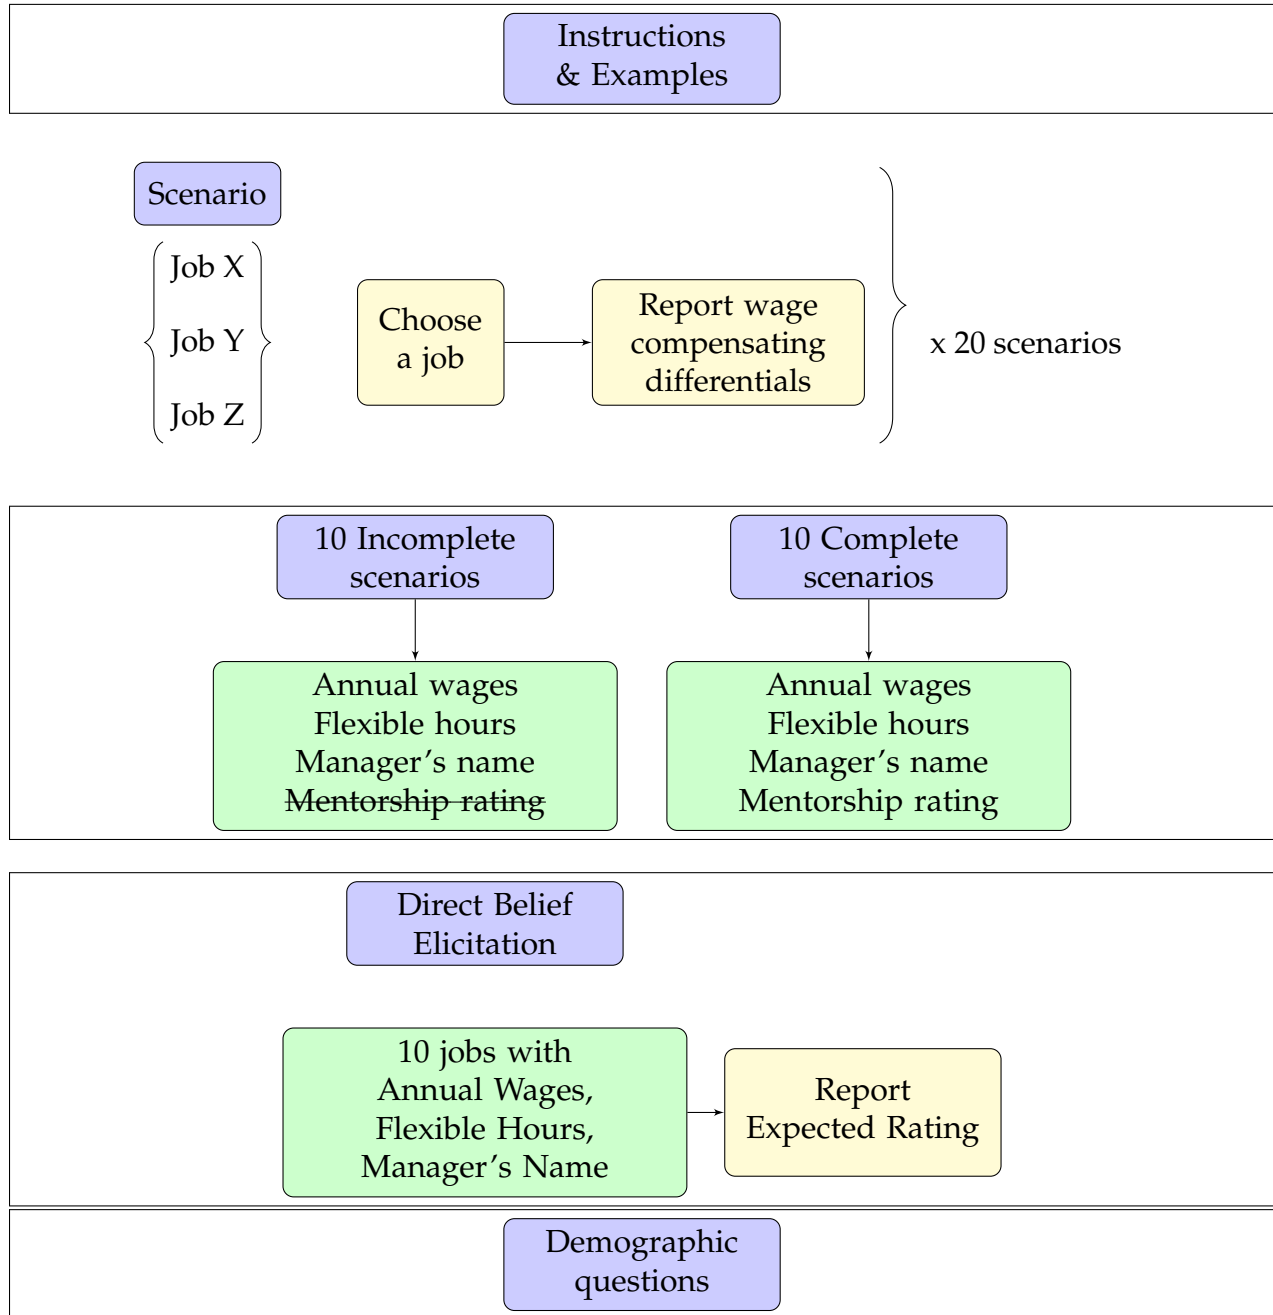
- Falk, A., Kosse, F. & Pinger, P. (2020), Mentoring and schooling decisions: Causal evidence. CESifo Working Paper No. 8382.
- Fisman, R. & O'Neill, M. (2009), 'Gender differences in beliefs on the returns to effort evidence from the world values survey', *Journal of Human Resources* **44**(4), 858–870.
- Flabbi, L., Macis, M., Moro, A. & Schivardi, F. (2019), 'Do female executives make a difference? the impact of female leadership on gender gaps and firm performance', *The Economic Journal* **129**(622), 2390–2423.
- Flory, J. A., Leibbrandt, A. & List, J. A. (2015), 'Do competitive workplaces deter female workers? a large-scale natural field experiment on job entry decisions', *The Review of Economic Studies* **82**(1), 122–155.
- Frederiksen, A., Kahn, L. B. & Lange, F. (2020), 'Supervisors and performance management systems', *Journal of Political Economy* **128**(6), 2123–2187.
- Fuster, A., Kaplan, G. & Zafar, B. (2021), 'What would you do with \$500? spending responses to gains, losses, news, and loans', *The Review of Economic Studies* **88**(4), 1760–1795.
- Guryan, J. & Charles, K. K. (2013), 'Taste-based or statistical discrimination: The economics of discrimination returns to its roots', *The Economic Journal* **123**(572), F417–F432.
- He, A. X. & le Maire, D. (2020), 'Mergers and managers: Manager-specific wage premiums and rent extraction in m&as', *Available at SSRN 3481262*.
- Heckman, J. J. (1998), 'Detecting discrimination', *Journal of Economic Perspectives* **12**(2), 101–116.
- Hoffman, M. & Tadelis, S. (2021), 'People management skills, employee attrition, and manager rewards: An empirical analysis', *Journal of Political Economy* **129**(1), 243–285.
- Hvidberg, K. B., Kreiner, C. & Stantcheva, S. (2020), Social position and fairness views, Technical report, National Bureau of Economic Research.
- Jäger, S., Roth, C., Roussille, N. & Schoefer, B. (2021), 'Worker beliefs about outside options and rents'.
- Kahn, L. B. (2013), 'Asymmetric information between employers', *American Economic Journal: Applied Economics* **5**(4), 165–205.

- Kline, P. M., Rose, E. K. & Walters, C. R. (2021), Systemic discrimination among large us employers, Technical report, National Bureau of Economic Research.
- Koşar, G., Ransom, T. & Van der Klaauw, W. (2021), ‘Understanding migration aversion using elicited counterfactual choice probabilities’, *Journal of Econometrics* .
- Koşar, G., Şahin, A. & Zafar, B. (2021), The work-leisure tradeoff: Identifying the heterogeneity, Working paper.
- Kuziemko, I., Norton, M. I., Saez, E. & Stantcheva, S. (2015), ‘How elastic are preferences for redistribution? evidence from randomized survey experiments’, *American Economic Review* **105**(4), 1478–1508.
- Lang, K. & Lehmann, J.-Y. K. (2012), ‘Racial discrimination in the labor market: Theory and empirics’, *Journal of Economic Literature* **50**(4), 959–1006.
- Lange, F. (2007), ‘The speed of employer learning’, *Journal of Labor Economics* **25**(1), 1–35.
- Lazear, E. P., Shaw, K. L. & Stanton, C. T. (2015), ‘The value of bosses’, *Journal of Labor Economics* **33**(4), 823–861.
- Lyle, D. S. & Smith, J. Z. (2014), ‘The effect of high-performing mentors on junior officer promotion in the us army’, *Journal of Labor Economics* **32**(2), 229–258.
- Mas, A. & Pallais, A. (2017), ‘Valuing alternative work arrangements’, *American Economic Review* **107**(12), 3722–59.
- Müller-Itten, M. & Öry, A. (2022), ‘Mentoring and the dynamics of affirmative action’, *American Economic Journal: Public Policy*,(forthcoming) .
- Parker, J. A. & Souleles, N. S. (2019), ‘Reported effects versus revealed-preference estimates: Evidence from the propensity to spend tax rebates’, *American Economic Review: Insights* **1**(3), 273–90.
- Phelps, E. S. (1972), ‘The statistical theory of racism and sexism’, *The American Economic Review* **62**(4), 659–661.
- Pinkston, J. C. (2009), ‘A model of asymmetric employer learning with testable implications’, *The Review of Economic Studies* **76**(1), 367–394.
- Robinson, J. (1933), ‘The economics of imperfect competition’.

- Rosen, S. (1986), The theory of equalizing differences, in R. l. Orley C. Ashenfelter, ed., 'Handbook of Labor Economics', Vol. 1, Elsevier, pp. 641–692.
- Sorkin, I. (2018), 'Ranking firms using revealed preference', *The quarterly journal of economics* **133**(3), 1331–1393.
- Stantcheva, S. (2021), Perceptions and preferences for redistribution, Working Paper 29370, National Bureau of Economic Research.
URL: <http://www.nber.org/papers/w29370>
- Stantcheva, S. (2022), How to run surveys: A guide to creating your own identifying variation and revealing the invisible, Working Paper 30527, National Bureau of Economic Research.
URL: <http://www.nber.org/papers/w30527>
- Taber, C. & Vejlín, R. (2020), 'Estimation of a roy/search/compensating differential model of the labor market', *Econometrica* **88**(3), 1031–1069.
- Wiswall, M. & Zafar, B. (2015), 'Determinants of college major choice: Identification using an information experiment', *The Review of Economic Studies* **82**(2), 791–824.
- Wiswall, M. & Zafar, B. (2018), 'Preference for the workplace, investment in human capital, and gender', *The Quarterly Journal of Economics* **133**(1), 457–507.

Tables and Figures

Figure 1: Schematic representation of survey flow



Notes: Every scenario consists of three different jobs: X, Y and Z. Individuals choose their most preferred job. For the jobs that they do not choose, individuals are asked to report the minimum increase in wages that they would need to choose those jobs instead. There are 20 such scenarios. In the first 10 scenarios, individuals do not observe the manager mentorship rating; however, in the last 10, they do, along with the other attributes.

Figure 2: Definitions of attributes

INSTRUCTIONS

In each question you will see a choice scenario.

A scenario will have 3 jobs (X, Y and Z), which you have to assume have been offered to you.

Each job will have 4 characteristics:

Manager: First name of the team's manager.

Annual wages: Gross annual salary (in lakhs).

Flexible hours: Whether the job allows for flexible hours or not.

Manager rating: This is the average rating of the mentorship of the manger, provided by this manager's current employees in an anonymous survey.

This is a measure of how good of a mentor is this manager to its subordinates.

This rating is on a scale of 1-5. The scale points mean as follows: 1: Poor; 2: Fair 3: Good, 4: Very good and 5: Excellent.

There will be 20 such choice scenarios.

Figure 3: Instructions

INSTRUCTIONS

In each scenario we will first ask you:

A. To choose one job among 3 job options.

Your instructions are:

To assume that all other characteristics, which are NOT MENTIONED here, are THE SAME in all the jobs. For example work from home under each manager is not shown here. You are to assume that its either available or not available in all three jobs. No job is different in anything that you do not observe,

In each scenario, we will then ask you for:

B. If you were to negotiate wages, how much **minimum increase in wages** you would need in each of the other two jobs, for you choose them instead.

Your instructions are:

State your wage increase, assuming that it does not change anything else about the job.

Let us give you an example from a survey of food preference.

A. Here you will see me choose one dish among 3 different dishes.

B. Then you will see me state how much MINIMUM drop in prices in the other two dishes would make me choose each of them instead.

Table 1: Incomplete scenario example

Job choice

Compensating differential (*if job chosen was Y*)

Rating of each manager could be different, but the data is unavailable. Anything else that you don't see here, is the SAME across all jobs. Please select your most preferred job.

Attributes	JOB X	JOB Y	JOB Z
Annual Wages	6	6.6	6.4
Flexible hours	yes	no	yes
Manager	Anirban	Shrinita	Arup
Manager rating	N/A	N/A	N/A

Job X

Job Y


Job Z

You chose Job Y.

If you were to negotiate your wage, how much of a MINIMUM INCREASE IN WAGES would you need in each of the other jobs for you to choose it instead of Job Y?

The scale here ranges from 0 to 2 lakhs.

Attributes	JOB X	JOB Y	JOB Z
Annual Wages	6	6.6	6.4
Flexible hours	yes	no	yes
Manager	Anirban	Shrinita	Arup
Manager rating	N/A	N/A	N/A



Notes: Jobs X and Z have male managers and Job Y has a female manager. Across the 10 incomplete scenarios, five scenarios have two jobs with male managers and one job with a female manager, and the remaining five have two jobs with female managers and one job with a male manager.

Table 2: Complete scenario example

Job choice

Anything else that you don't see here, is the SAME across all jobs. Manager rating is on a scale of 1-5. 1: Poor; 2: Fair; 3: Good; 4: Very good; 5: Excellent. Please select your most preferred job.

Attributes	JOB X	JOB Y	JOB Z
Annual Wages	8.2	8	8.6
Flexible hours	yes	no	yes
Manager	Mohan	Mohit	Mahima
Manager rating	3.70	4.00	3.15

Job X

Job Y

Job Z


Compensating differential example (if job chosen was X)

You chose Job X.

If you were to negotiate your wage, how much of a MINIMUM INCREASE IN WAGES would you need in each of the other jobs for you to choose it instead of Job X?

The scale here ranges from 0 to 2 lakhs.

Attributes	JOB X	JOB Y	JOB Z
Annual Wages	8.2	8	8.6
Flexible hours	yes	no	yes
Manager	Mohan	Mohit	Mahima
Manager rating	3.70	4.00	3.15



Notes: Jobs X and Y have male managers, and Job Z has a female manager. Overall, across the 10 complete scenarios, five scenarios have two jobs with male managers and one job with a female manager, and the remaining five have two jobs with female managers and one job with a male manager.

Table 3: Incomplete scenario adapted example for representative US jobs

Job choice

Compensating differential (if job chosen was Y)

Rating of each manager could be different, but the data is unavailable. Anything else that you don't see here, is the SAME across all jobs. Please select your most preferred job.

Attributes	JOB X	JOB Y	JOB Z
Annual Wages	35k	39k	37k
Flexible hours	yes	no	yes
Manager	John	Susan	Robert
Manager rating	N/A	N/A	N/A


Job X

Job Y

Job Z

You chose Job Y.
 If you were to negotiate your wage, how much of a MINIMUM INCREASE IN WAGES would you need in each of the other jobs for you to choose it instead of Job Y?
 The scale here ranges from 0 to 5,000 USD.

Attributes	JOB X	JOB Y	JOB Z
Annual Wages	35k	39k	37k
Flexible hours	yes	no	yes
Manager	John	Susan	Robert
Manager rating	N/A	N/A	N/A



Notes: Jobs X and Z have male managers, and Job Y has a female manager. Across the 10 incomplete scenarios, five scenarios have two jobs with male managers and one job with a female manager, and the remaining five have two jobs with female managers and one job with a male manager.

Table 4: Complete scenario adapted example for representative US jobs

Job choice

Anything else that you don't see here, is the SAME across all jobs. Manager rating is on a scale of 1-5. 1: Poor; 2: Fair; 3: Good; 4: Very good; 5: Excellent. Please select your most preferred job.

Attributes	JOB X	JOB Y	JOB Z
Annual Wages	40k	42k	45k
Flexible hours	yes	no	yes
Manager	James	Barbara	Mary
Manager rating	3.70	4.00	3.15

Job X

Job Y

Job Z


Compensating differential (if job chosen was X)

You chose Job X.

If you were to negotiate your wage, how much of a MINIMUM INCREASE IN WAGES would you need in each of the other jobs for you to choose it instead of Job X?

The scale here ranges from 0 to 5,000 USD.

Attributes	JOB X	JOB Y	JOB Z
Annual Wages	40k	42k	45k
Flexible hours	yes	no	yes
Manager	James	Barbara	Mary
Manager rating	3.70	4.00	3.15




Notes: Jobs X and Y have male managers, and Job Z has a female manager. Overall, across the 10 complete scenarios, five scenarios have two jobs with male managers and one job with a female manager, and the remaining five have two jobs with female managers and one job with a male manager.

Table 5: Summary statistics of job attributes

Variable	Mean	Std. Dev.	N
Female Manager	0.50	0.504	60
Flexible hours	0.53	0.503	60
Annual wages	7.11	1.476	60
Rating	3.41	0.495	30

Figure 4: Reported compensating differentials in unchosen jobs



Notes: The increase in wages is in units of 1 lakh (hundred thousand) INR. The figure is plotted for values only between 0 and 2 lakhs.

Table 6: Job attributes across male and female managers

Attribute	Male Manager	Female Manager	Difference	Std. Diff.
Rating	3.373 (0.519)	3.447 (0.485)	0.073 (0.183)	0.103
Flexible hours	0.567 (0.504)	0.500 (0.509)	-0.067 (0.131)	-0.093
Annual wages	7.080 (1.529)	7.140 (1.445)	0.060 (0.384)	0.029

Table 7: Sample demographics

	Gender		Total
	Male	Female	
580 respondents	58.3	41.7	100
Area of Study:			
Arts	25.1	71.9	44.7
Engineering	49.4	10.7	33.3
Science	25.4	17.4	22.1
Family Income:			
Less than 2 lakhs (Less than \$9,492)	37.6	24.4	32.1
2 lakhs to 5 lakhs (\$9,492 to \$23,730)	26.9	24.0	25.7
5 lakhs to 10 lakhs (\$23,730 to \$47,460)	21.3	30.2	25.0
10 to 20 lakhs (\$47,460 to \$94,921)	11.5	15.3	13.1
Above 20 lakhs (Above \$94,921)	2.7	6.2	4.1
Mother's Education:			
Below High School	17.8	11.2	15.0
High School	32.8	18.6	26.9
Bachelor's	38.5	46.3	41.7
Master's	8.0	16.9	11.7
Above Master's	3.0	7.0	4.7
Father's Education:			
Below High School	9.2	7.4	8.4
High School	21.3	11.2	17.1
Bachelor's	51.5	55.4	53.1
Master's	13.9	17.8	15.5
Above Master's	4.1	8.3	5.9
Mother's Occupation:			
Government	10.9	15.3	12.8
Homemaker	70.1	63.2	67.2
Not Applicable	4.4	3.3	4.0
Private Sector	4.7	8.7	6.4
Self-Employed	9.8	9.5	9.7
Father's Occupation:			
Government	30.8	33.1	31.7
Homemaker	3.3	0.8	2.2
Not Applicable	16.3	11.2	14.1
Private Sector	16.9	20.7	18.4
Self-Employed	32.8	34.3	33.4

Notes: All variables are categorical. Numbers represent percentages. The parental occupation category of "Not Applicable" refers to a deceased parent. Variables on income categories are in INR and have their corresponding purchasing power parity-adjusted USD equivalent below each category.

Table 8: Percentages of chosen jobs with male and female managers, by scenario

<i>Respondent</i>	<i>All</i>		<i>Female</i>		<i>Male</i>	
<i>Scenario</i> \ Manager	Male	Female	Male	Female	Male	Female
Incomplete	48.2	51.8	48.1	52.9	48.3	51.7
Complete	38.9	61.1	38.8	61.2	38.9	61.1

Table 9: Average compensating differentials for unchosen jobs, by manager gender across scenarios

	Male Manager	Female Manager	Difference	St Diff
Incomplete scenarios	0.957 (0.965)	1.020 (0.998)	-0.063*** (0.018)	-0.046
Complete scenarios	1.100 (1.058)	1.038 (0.977)	0.061*** (0.019)	0.042

Notes: Units are in 1 lakh (hundred thousand) INR

Table 10: Difference-in-differences estimates from the complete scenarios data

VARIABLES	Margins
Female Worker	-0.001 (0.010)
Female Manager	0.081*** (0.009)
Female Worker X Female Manager	0.001 (0.014)
Annual Wages	0.347*** (0.019)
Mentorship Rating	0.484*** (0.009)
Flexible Hours	0.262*** (0.009)
Scenario FE	yes
Observations	17,400

Notes: The estimates show the marginal effects of each of the attributes in a difference-in-differences specification. Standard errors bootstrapped at the individual level with 1,000 replications. The total number of observations is 17,400 because we use individual-level choice data on 3 jobs in each of the 10 complete scenarios for 580 individuals.

Table 11: Complete and incomplete scenarios: Jointly estimated

Incomplete Scenarios			Complete Scenarios		
<i>Parameters</i>	in 10 ⁵ INR	% of wages	<i>Parameters</i>	in 10 ⁵ INR	% of wages
$\tilde{\beta}_i^G = \beta_i^G + \beta_i^R \alpha_i^G$	-0.007 (0.012)	-0.01%	β_i^G (Male Manager)	-0.118*** (0.019)	-1.7%
$\tilde{\beta}_i^H = \beta_i^H + \beta_i^R \alpha_i^H$	1.136*** (0.068)	16.1%	β_i^H (Flexible Hours)	0.776*** (0.027)	11.1%
$\tilde{\beta}_i^W = \beta_i^W + \beta_i^R \alpha_i^W$	1.132*** (0.059)	16%	β_i^W (Annual Wages)	1	
			β_i^R (Mentorship)	0.793*** (0.029)	11.3%
Observations	11,600			11,600	

Notes: The table shows estimates from estimating equation system (14) for each individual and reports the averages $\mathbb{E}(\tilde{\beta}_i)$ in the incomplete scenarios and $\mathbb{E}(\beta_i)$ in the complete scenarios. β_i^W is normalized to 1. Estimates are represented in two sets of units—the first is in hundred thousand INR, and the second converts those units into percentages of average annual wages. Standard errors are computed using the the block bootstrap at the individual level with 1,000 replications. Statistical significance at 1, 5, and 10% is denoted by ***, **, and *, respectively.

Table 12: Estimates of belief parameters

Belief Parameter	
α_i^G (Male managers)	0.140*** (0.024)


Notes: Standard errors are computed using the block bootstrap at the individual level with 1,000 repetitions. Parameters are estimated by jointly estimating the complete and incomplete scenarios and using the estimates of β^R , $\tilde{\beta}^G$ and β^G presented in Table 11. The estimates come from the equation $\alpha_i^G = \frac{\tilde{\beta}_i^G - \beta_i^G}{\beta_i^R}$. Statistical significance at 1, 5, and 10% is denoted by ***, **, and *, respectively.

Table 13: Estimates of valuation of beliefs on male relative to female manager mentorship ability

Belief Parameter	in 10^5 INR	% of wages
$\beta_i^R \alpha_i^G$ (Male managers)	0.112*** (0.024)	1.6%


Notes: Estimates come from the equation $\beta^R \alpha^G = \tilde{\beta}^G - \beta^G$ for all $x \in \{G, W, H\}$ presented in Table 11. The standard errors are block bootstrapped at the individual level with 1,000 repetitions. Parameters are estimated by jointly estimating the complete and incomplete scenarios. Statistical significance at 1, 5, and 10% is denoted by ***, **, and *, respectively. Average annual wages equal 7 lakh INR (\approx \$38.8 thousand in PPP).

Figure 5: Distribution of valuation of individual beliefs as a percentage of average annual wages



Notes: The figure plots the histogram and the smoothed empirical cumulative distribution function of individual beliefs obtained by estimating the model for each individual. The unit on the x-axis is the percentage of average annual wages.

Figure 6: Distribution of valuation of individual preferences as a percentage of average annual wages



Notes: The figure plots the histogram and the smoothed empirical cumulative distribution function of individual preferences obtained by estimating the model for each individual. The unit on the x-axis is the percentage of average annual wages.

Table 14: Correlations of estimated parameters with individual characteristics

	Preferences	Valuation of beliefs
Female	-0.000 (0.023)	-0.010 (0.037)
<i>Major</i>		
Engineering	-0.063** (0.032)	0.044 (0.052)
Science	-0.039 (0.027)	0.023 (0.042)
<i>Education</i>		
Masters	-0.018 (0.029)	0.052 (0.048)
MPhil/PhD	0.052 (0.080)	-0.045 (0.124)
<i>Family income</i>		
2 to 5 lakhs (\$9,492 to \$23,730)	0.009 (0.027)	-0.010 (0.043)
5 to 10 lakhs (\$23,730 to \$47,460)	-0.002 (0.028)	-0.027 (0.045)
10 to 20 lakhs (\$47,460 to \$94,921)	0.008 (0.034)	0.014 (0.054)
Above 20 lakhs (Above \$94,921)	-0.035 (0.054)	0.040 (0.090)
1(Mother education \geq father education)	0.007 (0.060)	-0.194** (0.093)
1(Father employed)	-0.049 (0.041)	0.198*** (0.065)
1(Mother employed)	0.019 (0.030)	-0.002 (0.049)
<i>Mother more educated than father \times Father employed \times Mother employed</i>		
NNY	0.041 (0.093)	-0.170 (0.144)
NYN	0.039 (0.065)	-0.218** (0.102)
NYY	0.041 (0.075)	-0.273** (0.119)
YNY	-0.043 (0.069)	0.088 (0.111)
Observations	578	535
R^2	0.01	0.01

Notes: Error terms in both regressions are homoskedastic. This is because the underlying identification of the preference and belief parameters is within-individual, which leads to the error terms being i.i.d.

Figure 7: Expected rating

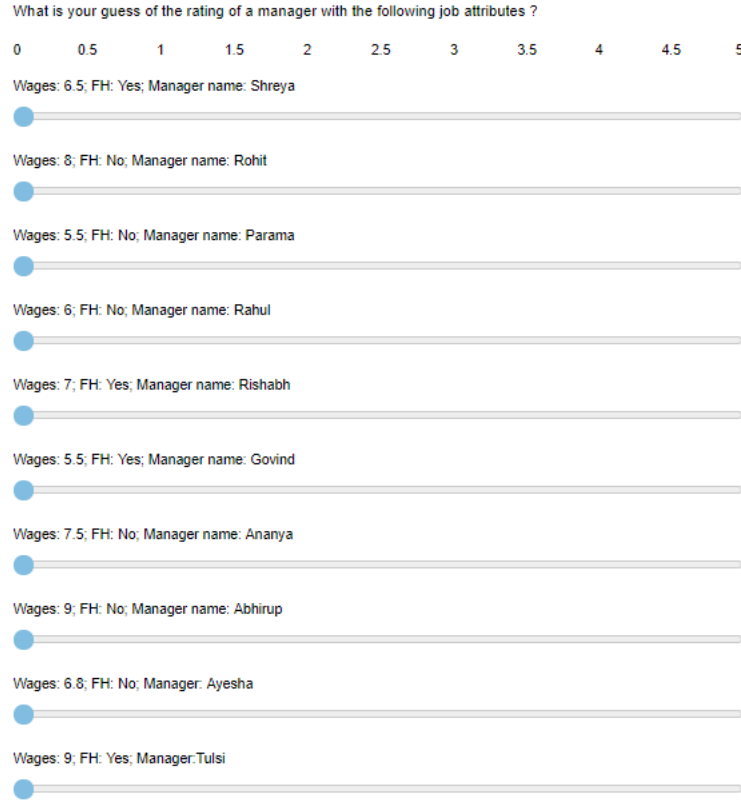


Table 15: Estimates of expected beliefs from directly elicited belief data

Belief Parameters	Model			Data on beliefs		
	All	Women	Men	All	Women	Men
α^G (Male managers)	0.140*** (0.024)	0.108*** (0.033)	0.161*** (0.033)	0.091*** (0.022)	0.099*** (0.034)	0.084*** (0.033)
α^W (Annual wages)	0.465*** (0.053)	0.478*** (0.077)	0.457*** (0.074)	0.428*** (0.003)	0.425*** (0.004)	0.429*** (0.004)
α^H (Flexible hours)	0.449*** (0.040)	0.422*** (0.060)	0.467*** (0.056)	0.358*** (0.021)	0.352*** (0.032)	0.362*** (0.030)

Notes: Standard errors bootstrapped at the individual level with 1,000 repetitions. These estimates come from projecting a linear model on the data of the reported expected manager rating in each of 10 different jobs with annual wages, flexibility of hours and manager name. The linear projection takes the form of $R_{ij} = \alpha^G G_j + \alpha^W W_j + \alpha^H H_j + \eta_{ij}$.

A Appendix

A.1 Identification

Let individual $i \in \{1, \dots, N\}$ have preferences on attributes X_j in job $j \in \{1, \dots, J\}$ given by the function

$$U_{ij} = u_i(X_j) + \epsilon_{ij}$$

Identification is achieved in two steps. We first show the identification of preferences when information is complete in Step 1. Then, with incomplete information in Step 2, we show that we can identify beliefs given preferences.

Step 1: Identifying preferences

Individuals form expectations when they report their job choices and compensating differentials. Hence,

$$\mathbb{E}_i[U_{ij}|X_j] = u_i(X_j) + \mathbb{E}_i[\epsilon_{ij}|X_j]$$

A compensating differential of Δ_{ijk} makes i indifferent between jobs j and k when i observes X_j and X_k . Hence, given that a compensating differential increases only wages and changes nothing else about the job, we have $\mathbb{E}_i[\epsilon_{ik}|X_k, \Delta_{ijk}] = \mathbb{E}_i[\epsilon_{ik}|X_k]$. Normalizing the preference parameter on wages to 1, we have

$$\begin{aligned} \mathbb{E}_i[U_{ik}|X_k, \Delta_{ijk}] &= \Delta_{ijk} + u_i(X_k) + \mathbb{E}_i[\epsilon_{ik}|X_k, \Delta_{ijk}] \\ &= \Delta_{ijk} + u_i(X_k) + \mathbb{E}_i[\epsilon_{ik}|X_k] \end{aligned}$$

Given that everything else about the job is the same, we have $\mathbb{E}_i[\epsilon_{ik}|X_k] = \mathbb{E}_i[\epsilon_{ij}|X_j]$ for all jobs $j \neq k$.

Since by definition $\Delta_{ijk} = \mathbb{E}_i[U_{ij}|X_j] - \mathbb{E}_i[U_{ik}|X_k, \Delta_{ijk}]$, we now have

$$\Delta_{ijk} = u_i(X_j) - u_i(X_k)$$

As the number of scenarios goes to infinity, for each individual i , this identifies preferences $u_i(\cdot)$ as long as $\text{Var}_i(u_i(X_j)|X_j) \neq 0$.

Step 2: Identifying beliefs given preferences

Now, when i observes $\tilde{X}_j \equiv X_j \setminus R_j$ for each job $j \in \{1, \dots, J\}$, i forms beliefs on R_j given \tilde{X}_j according to $G_i(R|\tilde{X})$. Now, we have

$$\mathbb{E}_i[U_{ij}|\widetilde{X}_j] = \mathbb{E}_i[u_i(X_j)|\widetilde{X}_j] + \mathbb{E}_i[\epsilon_{ij}|\widetilde{X}_j]$$

A compensating differential of $\widetilde{\Delta}_{ijk}$ makes i indifferent between jobs j and k while observing \widetilde{X}_j and \widetilde{X}_k , respectively. Hence, given that the compensating differential increases only wages and changes nothing else about the job, and normalizing the preference parameter on wages to 1, we have

$$\begin{aligned}\mathbb{E}_i[U_{ik}|\widetilde{X}_k, \widetilde{\Delta}_{ijk}] &= \widetilde{\Delta}_{ijk} + \mathbb{E}_i[u_i(X_k)|\widetilde{X}_k] + \mathbb{E}_i[\epsilon_{ik}|\widetilde{X}_k, \widetilde{\Delta}_{ijk}] \\ &= \widetilde{\Delta}_{ijk} + \mathbb{E}_i[u_i(X_k)|\widetilde{X}_k] + \mathbb{E}_i[\epsilon_{ik}|\widetilde{X}_k]\end{aligned}$$

Given that everything else about the job is the same, we have $\mathbb{E}_i[\epsilon_{ik}|\widetilde{X}_k] = \mathbb{E}_i[\epsilon_{ij}|\widetilde{X}_j]$ for all jobs $j \neq k$.

Since by definition $\widetilde{\Delta}_{ijk} = \mathbb{E}_i[U_{ij}|\widetilde{X}_j] - \mathbb{E}_i[U_{ik}|\widetilde{X}_k, \widetilde{\Delta}_{ijk}]$, we have

$$\begin{aligned}\widetilde{\Delta}_{ijk} &= \mathbb{E}_i[u_i(X_j)|\widetilde{X}_j] - \mathbb{E}_i[u_i(X_k)|\widetilde{X}_k] \\ &= \int u_i(X_j) dG_i(R|\widetilde{X} = \widetilde{X}_j) - \int u_i(X_k) dG_i(R|\widetilde{X} = \widetilde{X}_k)\end{aligned}$$

Assuming that moments exist, as the number of scenarios go to infinity, the above identifies individual i 's belief distribution $G_i(R|\widetilde{X})$, given $\widetilde{\Delta}_{ijk}$ and $u_i(\cdot)$ identified from Step 1. The model is also identified as the number of attributes goes to infinity as long as it approaches infinity at a slower rate than the number of scenarios approaches infinity.

Note that beliefs $G_i(R|\widetilde{X})$ are not identified under two circumstances. First, if i does not care about R , then the choices made and the compensations reported are not driven by whatever way i may expect R to vary with \widetilde{X} . To see this mathematically, if i does not care about R , then the above set of equations become independent of $G_i(R|\widetilde{X})$ because $\mathbb{E}_i[u_i(X_j)|\widetilde{X}_j] = \mathbb{E}_i[u_i(\widetilde{X}_j)|\widetilde{X}_j] = u_i(\widetilde{X}_j)$. The first equality follows from $\widetilde{X}_j \equiv X_j \setminus R_j$, and i does not care about R . Second, if R is independent of \widetilde{X} , then no variation in \widetilde{X} can generate any variation in the beliefs and thus will not be reflected in the choices and compensating differentials. To see this mathematically, if R is independent of \widetilde{X} , then $G_i(R|\widetilde{X}) = G_i(R)$. This makes $\mathbb{E}_i[u_i(X_j)|\widetilde{X}_j] \equiv \mathbb{E}_i[u_i(\widetilde{X}_j, R_j)|\widetilde{X}_j]$ a function that is independent of R_j by the law of iterated expectations.

A.2 Estimation details

We implement the joint estimation of individual indifference in the complete and incomplete scenarios in a fully interacted model by stacking up the matrices of observables across the complete and incomplete scenarios. In particular, we estimate the following constrained least squares regression with no constant:

$$\begin{bmatrix} \Delta_{jk} \\ \vdots \\ \tilde{\Delta}_{jk} \\ \vdots \end{bmatrix} = \begin{bmatrix} \mathbf{X}_j - \mathbf{X}_k & \mathbf{0} \\ \vdots & \vdots \\ \mathbf{0} & \tilde{\mathbf{X}}_j - \tilde{\mathbf{X}}_k \\ \vdots & \vdots \end{bmatrix}' (\boldsymbol{\beta} \quad \tilde{\boldsymbol{\beta}}) + \mathbf{e} \quad (16)$$

where the constraint is the normalization for the preference parameter on wages to be equal to one. The standard errors are computed using the block bootstrap at the student level. This accounts for any arbitrary correlation between responses at the student level.

The block bootstrap algorithm is as follows: The sample contains N individuals.

1. Generate B the block bootstrap samples of N individuals each.
For each $b = 1, \dots, B$
2. Estimate the model for each member in the bootstrap sample by bootstrapping each member's responses.

Obtain $\hat{\beta}_i^{(b)}$ and compute its sample mean $\hat{\beta}^{(b)} = \sum_{i=1}^N \hat{\beta}_i^{(b)}$.

Compute the mean and standard deviation of the B estimates in hand to generate estimates of the bootstrap mean and bootstrap standard error.

Preferences:

$$\begin{aligned} \hat{\mathbb{E}}(\beta_i) &= \frac{1}{B} \sum_{b=1}^B \hat{\beta}^{(b)} \\ \widehat{\text{std error}}(\beta_i) &= SD(\hat{\beta}^{(b)}) \end{aligned}$$

Preferences confounded with valuation of beliefs:


$$\begin{aligned} \hat{\mathbb{E}}(\tilde{\beta}_i) &= \frac{1}{B} \sum_{b=1}^B \hat{\tilde{\beta}}^{(b)} \\ \widehat{\text{std error}}(\tilde{\beta}_i) &= SD(\hat{\tilde{\beta}}^{(b)}) \end{aligned}$$

Valuation of beliefs:

$$\hat{\mathbb{E}}(\tilde{\beta}_i - \beta_i) = \frac{1}{B} \sum_{b=1}^B \widehat{\tilde{\beta} - \beta}^{(b)}$$

$$\widehat{\text{std error}}(\tilde{\beta}_i - \beta_i) = SD(\widehat{\tilde{\beta} - \beta}^{(b)})$$

Figure 8: Bootstrap distribution of beliefs and preferences on male managers



Notes: The figure shows bootstrap distributions of beliefs and preferences in the percentage of average annual wages. The bootstrap distributions are obtained from 1,000 block bootstrapped samples, using the algorithm described in Appendix A.2. These bootstrap distributions are used for estimation of means and standard errors of preferences and beliefs.

A.3 Alternate models

In this section, we illustrate some examples by relaxing the linearly separable model to include interactions. These models are identified as shown in the Appendix above.

Each individual $i \in \{1, \dots, N\}$ has preferences on attributes X_j in job $j \in \{1, \dots, J\}$ given by the function

$$U_{ij} = u_i(X_j) + \epsilon_{ij}$$

Individuals observe $X_j \equiv \{G_j, W_j, H_j, R_j\}$ in the complete scenarios and $\tilde{X}_j \equiv \{G_j, W_j, H_j\}$ in the incomplete scenarios.

Throughout this section, we maintain the same assumptions parallel to our instructions:

Assumption (1): All attributes not mentioned in the survey are the same for all jobs.

Assumption (2): The reported compensating differential increases only wages and changes nothing else about the job.

We specify the belief function as

$$\mathbb{E}(R_j | \tilde{X}_j) = \sum_{x \in \tilde{X}} \alpha_i^x x_j$$

Note that in these models, since we are interacting the mentorship rating with other observables, for the model to be identified, we have to assume $\mathbb{E}_i(\eta_{ij} | \tilde{X}_j) = 0$ in the belief function. This was not required when the mentorship rating appeared linearly separable in the model.

A.3.1 Model with an interaction of manager gender (G) and manager mentorship rating (R)

In this example, the utility of individual i is given by

$$U_{ij} = \sum_{x \in X} \beta_i^x x_j + \beta_i^{GR} G_j R_j + \epsilon_{ij}$$

Thus, the parameter space now contains 5 preference parameters:

$$\beta_i \equiv \{\beta_i^G, \beta_i^H, \beta_i^W, \beta_i^R, \beta_i^{GR}\}$$

and 4 belief parameters:

$$\alpha_i \equiv \{\alpha_i^G, \alpha_i^H, \alpha_i^W\}$$

Complete scenarios

In the complete scenarios, we have for all i

$$\mathbb{E}_i[U_{ij}|X_j] = \sum_{x \in X} \beta_i^x x_j + \beta_i^{GR} G_j R_j + \mathbb{E}_i(\epsilon_{ij}|X_j)$$

In any other job k that is not chosen, supposing that the individual reports Δ_{ijk} as the compensating differential, we have

$$\mathbb{E}_i[U_{ik}|X_k, \Delta_{ijk}] = \sum_{x \in X} \beta_i^x x_k + \beta_i^{GR} G_k R_k + \beta_i^W \Delta_{ijk} + \mathbb{E}_i(\epsilon_{ik}|X_k, \Delta_{ijk})$$

Given assumption (2), we have $\mathbb{E}_i(\epsilon_{ik}|X_k, \Delta_{ijk}) = \mathbb{E}_i(\epsilon_{ik}|X_k)$, which by assumption (1) is equal to $\mathbb{E}_i(\epsilon_{ij}|X_j)$.

Thus, normalizing $\beta_i^W = 1$, we have,

$$\Delta_{ijk} = \sum_{x \in X} \beta_i^x (x_j - x_k) + \beta_i^{GR} (G_j R_j - G_k R_k) \quad (17)$$

Incomplete scenarios

In the incomplete scenarios, we have

$$\mathbb{E}_i[U_{ij}|\tilde{X}_j] = \sum_{x \in \tilde{X}} \beta_i^x x_j + \mathbb{E}_i(\beta_i^R R_j + \beta_i^{GR} G_j R_j + \epsilon_{ij}|\tilde{X}_j)$$

Assuming as before that individuals know their preference parameters, this simplifies to

$$\mathbb{E}_i[U_{ij}|\tilde{X}_j] = \sum_{x \in \tilde{X}} \beta_i^x x_j + \beta_i^R \mathbb{E}_i(R_j|\tilde{X}_j) + \beta_i^{GR} G_j \mathbb{E}_i(R_j|\tilde{X}_j) + \mathbb{E}_i(\epsilon_{ij}|\tilde{X}_j)$$

Using the belief function specified above, we have,

$$\mathbb{E}_i[U_{ij}|\tilde{X}_j] = \sum_{x \in \tilde{X}} (\beta_i^x + \beta_i^R \alpha_i^x) x_j + \beta_i^{GR} G_j \sum_{x \in \tilde{X}} \alpha_i^x x_j + \mathbb{E}_i(\epsilon_{ij}|\tilde{X}_j)$$

In any other job k that is not chosen, supposing that the individual reports $\tilde{\Delta}_{ijk}$ as the compensating differential, we have

$$\begin{aligned}\mathbb{E}_i[U_{ik}|\tilde{X}_k, \tilde{\Delta}_{ijk}] &= \sum_{x \in \tilde{X}} (\beta_i^x + \beta_i^R \alpha_i^x) x_k + \beta_i^{GR} G_k \sum_{x \in \tilde{X}} \alpha_i^x x_k + \beta_i^W \tilde{\Delta}_{ijk} \\ &+ \mathbb{E}_i(\epsilon_{ik}|\tilde{X}_k, \tilde{\Delta}_{ijk})\end{aligned}$$

Given assumption (2), we have $\mathbb{E}_i(\epsilon_{ik}|\tilde{X}_k, \tilde{\Delta}_{ijk}) = \mathbb{E}_i(\epsilon_{ik}|\tilde{X}_k)$, which by assumption (1) is equal to $\mathbb{E}_i(\epsilon_{ij}|\tilde{X}_j)$.

Normalizing $\beta_i^W = 1$, we have

$$\tilde{\Delta}_{ijk} = \sum_{x \in \tilde{X}} (\beta_i^x + \beta_i^R \alpha_i^x) (x_j - x_k) + \beta_i^{GR} \left(G_j \sum_{x \in \tilde{X}} \alpha_i^x x_j - G_k \sum_{x \in \tilde{X}} \alpha_i^x x_k \right)$$

which simplifies to

$$\tilde{\Delta}_{ijk} = \sum_{x \in \tilde{X}} (\beta_i^x + \beta_i^R \alpha_i^x) (x_j - x_k) + \beta_i^{GR} \sum_{x \in \tilde{X}} \alpha_i^x (x_j G_j - x_k G_k) \quad (18)$$

An important difference to note in this model relative to the linearly separable model is that β_i^G alone no longer captures how much worker i values a male manager over a female manager. In this case, it is

$$U_{ij}|_{G_j=1} - U_{ij}|_{G_j=0} = \beta_i^G + \beta_i^{GR} R_j$$

Thus, the average valuation of a male manager by worker i is $\beta_i^G + \beta_i^{GR} \mathbb{E}(R_j)$.³⁷

The valuation of beliefs is also different from that in the linearly separable model. This is because it is no longer weighted only by how much individuals care about mentorship quality (β_i^R) and now is augmented by how much individuals care about mentorship quality by manager gender. Using the equations above, we can show that the valuation of beliefs on mentorship for male managers relative to its counterpart for female managers in this model is

$$\alpha_i^G (\beta_i^R + \beta_i^{GR})$$

The incomplete scenarios identify $\beta_i^G + \alpha_i^G (\beta_i^R + \beta_i^{GR})$, $\beta_i^H + \beta_i^R \alpha_i^H$, $\beta_i^W + \beta_i^R \alpha_i^W$ and $\{\alpha_i^x \beta_i^{GR}\}_{x \in \{H, W\}}$. Given that the complete scenarios identify $\{\beta_i^x\}_{x \in X}$ and β_i^{GR} , the in-

³⁷The expectation does not condition on manager gender because the attributes are exogenously provided to the respondents and thus the mentorship rating does not significantly differ between male and female managers.

complete scenarios identify $\{\alpha_i^x\}_{x \in \tilde{X}}$.³⁸

A.3.2 Model with interaction of (G) and (H) in the beliefs for rating

We specify the belief function now as

$$\mathbb{E}_i(R_j | \tilde{X}_j) = \sum_{x \in \tilde{X}} \alpha_i^x x_j + \alpha_i^{GW} G_j H_j$$

The utility function is unchanged at

$$U_{ij} = \sum_{x \in X} \beta_i^x x_j + \epsilon_{ij}$$

From the complete scenarios, following similar steps, we can derive

$$\Delta_{ijk} = \sum_{x \in X} \beta_i^x (x_j - x_k) \quad (19)$$

From the incomplete scenarios, following similar steps, we can derive

$$\tilde{\Delta}_{ijk} = \sum_{x \in \tilde{X}} \left(\beta_i^x + \beta_i^R \alpha_i^x \right) (x_j - x_k) + \beta_i^R \alpha_i^{GW} (G_j H_j - G_k H_k) \quad (20)$$

Following similar arguments, we can show that all parameters in this model are identified.

In this model, the valuation of a male manager relative to that of a female manager by worker i is given as β_i^G . However, the valuation of beliefs will no longer be $\beta_i^R \alpha_i^G$. The valuation of beliefs of worker i on male managers' mentorship ability relative to female managers' mentorship ability in job j is

³⁸Similarly, if one were to use an interaction of the mentorship rating (R) with flexible hours (H), the corresponding complete scenario equation would be

$$\Delta_{ijk} = \sum_{x \in X} \beta_i^x (x_j - x_k) + \beta_i^{HR} (H_j R_j - H_k R_k)$$

, and the incomplete scenario equation would be

$$\tilde{\Delta}_{ijk} = \sum_{x \in \tilde{X}} \left(\beta_i^x + \beta_i^R \alpha_i^x \right) (x_j - x_k) + \sum_{x \in \tilde{X}} \alpha_i^x \beta_i^{HR} (x_j H_j - x_k H_k)$$

$$\begin{aligned} & \beta_i^R [\mathbb{E}_i(R_j|G_j = 1, W_j, H_j) - \mathbb{E}_i(R_j|G_j = 0, W_j, H_j)] \\ &= \beta_i^R [\alpha_i^G + \alpha_i^{GW} H_j] \end{aligned}$$

The valuation of the corresponding beliefs by worker i averaged over jobs is $\beta_i^R [\alpha_i^G + \alpha_i^{GW} \mathbb{E}(H_j)]$. Given the distribution of H_j and the identified preference and belief parameters for each worker i , the above equation identifies the distribution of the valuation of beliefs on male managers' mentorship rating relative to female managers'.

A.3.3 Model with log(wages)

In this model, the only minor difference is in the estimating equations since now the compensating differential is not separable from the wages due to the nonlinear logarithmic function. We keep other parts of the model linearly separable for simplicity; however, they can be relaxed as shown in the previous subsections.

$$U_{ij} = \beta_i^G G_j + \beta_i^H H_j + \beta_i^W \log(w_j) + \beta_i^R R_j + \epsilon_{ij}$$

Thus, the parameter space now contains 4 preference parameters:

$$\beta_i \equiv \{\beta_i^G, \beta_i^H, \beta_i^R\}$$

with β_i^W normalized to 1, and 4 belief parameters:

$$\alpha_i \equiv \{\alpha_i^G, \alpha_i^H, \alpha_i^W\}$$

Complete scenarios

In the complete scenarios, we have for all i

$$\mathbb{E}_i[U_{ij}|X_j] = \beta_i^G G_j + \beta_i^H H_j + \beta_i^W \log(w_j) + \beta_i^R R_j + \mathbb{E}_i(\epsilon_{ij}|X_j)$$

In any other job k that is not chosen, supposing that the individual reports Δ_{ijk} as the compensating differential, we have

$$\mathbb{E}_i[U_{ik}|X_k, \Delta_{ijk}] = \beta_i^G G_k + \beta_i^H H_k + \beta_i^W \log(w_k + \Delta_{ijk}) + \beta_i^R R_k + \mathbb{E}_i(\epsilon_{ik}|X_k, \Delta_{ijk})$$

Given assumption (2), we have $\mathbb{E}_i(\epsilon_{ik}|X_k, \Delta_{ijk}) = \mathbb{E}_i(\epsilon_{ik}|X_k)$, which by assumption (1)

is equal to $\mathbb{E}_i(\epsilon_{ij}|X_j)$.

Thus, normalizing $\beta_i^W = 1$, we have

$$\frac{\log(w_k + \Delta_{ijk})}{\log(w_j)} = \beta_i^G(G_j - G_k) + \beta_i^H(H_j - H_k) + \beta_i^R(R_j - R_k) \quad (21)$$

Incomplete scenarios

In the incomplete scenarios, we have

$$\mathbb{E}_i[U_{ij}|X_j] = \beta_i^G G_j + \beta_i^H H_j + \beta_i^W \log(w_j) + \beta^R \mathbb{E}_i(R_j|\tilde{X}_j) + \mathbb{E}_i(\epsilon_{ij}|\tilde{X}_j)$$

Using the belief function specified above, we have

$$\mathbb{E}_i[U_{ij}|\tilde{X}_j] = (\beta_i^G + \beta_i^R \alpha_i^G)G_j + (\beta_i^H + \beta_i^R \alpha_i^H)H_j + (\beta_i^W + \beta_i^R \alpha_i^W)\log(w_j) + \mathbb{E}_i(\epsilon_{ij}|\tilde{X}_j)$$

For job k and compensating differential $\tilde{\Delta}_{ijk}$, we have

$$\mathbb{E}_i[U_{ik}|\tilde{X}_k, \tilde{\Delta}_{ijk}] = (\beta_i^G + \beta_i^R \alpha_i^G)G_k + (\beta_i^H + \beta_i^R \alpha_i^H)H_k + \beta_i^W \log(w_k + \tilde{\Delta}_{ijk}) + \beta_i^R \alpha_i^W \log(w_k) + \mathbb{E}_i(\epsilon_{ij}|\tilde{X}_j)$$

This stems from the assumption that the compensating differential does not change anything about the job except for the wages. Hence, this will not change the expected rating. In particular,

$$\begin{aligned} \mathbb{E}(R_k|\tilde{X}_k, \tilde{\Delta}_{ijk}) &= \mathbb{E}(R_k|\tilde{X}_k) \\ &= \alpha_i^G G_k + \alpha_i^H H_k + \alpha_i^W \log(w_k) \end{aligned}$$

Thus, normalizing $\beta_i^W = 1$, we have

$$\frac{\log(w_k + \tilde{\Delta}_{ijk})}{\log(w_j)} = (\beta_i^G + \beta_i^R \alpha_i^G)(G_j - G_k) + (\beta_i^H + \beta_i^R \alpha_i^H)(H_j - H_k) + \beta_i^R \alpha_i^W (\log(w_j) - \log(w_k)) \quad (22)$$

A.4 Model with mentorship quality as a proxy for overall manager quality

In this section, we delineate a more generic model than the one presented in the main paper. The identifying assumptions remain the same. However, we relax the interpretation of the attribute of manager mentorship ability. In this specification, individuals care about overall manager quality (Q) in addition to caring about wages, flexibility of hours and manager gender. The mentorship rating acts as a signal of overall manager quality. Individuals care about this overall manager quality. The purpose of this generic model is to show that the finding of statistical discrimination still holds.

Redefining the set of attributes to $A \equiv \{G, W, H, Q\}$, the utility of individual i takes the same linear form:

$$U_{ij} = \sum_{x \in A} \beta_i^x x_j + \epsilon_{ij} \quad (23)$$

Observe that now in both the complete and incomplete scenarios, individuals need to form expectations on manager quality. In the incomplete scenarios, they do not have information on the manager's mentorship rating, whereas in the complete scenarios, they do. The expected utilities in the complete and incomplete scenarios take the following forms:

$$\begin{aligned} \text{Incomplete Scenarios: } \mathbb{E}_i[U_{ij}|\tilde{X}_j] &= \sum_{x \in A \setminus Q} \beta_i^x x_j + \beta_i^Q \mathbb{E}_i(Q_j|\tilde{X}_j) + \mathbb{E}_i(\epsilon_{ij}|\tilde{X}_j) \\ \text{Complete Scenarios: } \mathbb{E}_i[U_{ij}|X_j] &= \sum_{x \in A \setminus Q} \beta_i^x x_j + \beta_i^Q \mathbb{E}_i(Q_j|X_j) + \mathbb{E}_i(\epsilon_{ij}|X_j) \end{aligned} \quad (24)$$

As explained above, in both expected utilities, the individuals forms expectations on manager quality. However, in the complete scenario, the individual has the additional information of the manager's mentorship rating. We parameterize the expectation on manager quality in the following way:

$$\begin{aligned} \text{Complete Scenarios: } \mathbb{E}_i(Q_j|X_j) &= \sum_{x \in A \setminus Q} \gamma_i^x x_j + \gamma_i^R R_j + \mathbb{E}_i(\zeta_i|X_j) \\ \text{Incomplete Scenarios: } \mathbb{E}_i(Q_j|\tilde{X}_j) &= \sum_{x \in A \setminus Q} \gamma_i^x x_j + \gamma_i^R \mathbb{E}[R_j|\tilde{X}_j] + \mathbb{E}_i(\zeta_i|\tilde{X}_j) \\ &= \sum_{x \in A \setminus Q} \tilde{\gamma}_i^x x_j + \mathbb{E}_i(\zeta_i|\tilde{X}_j) \end{aligned} \quad (25)$$

Incorporating the above in the expected utility functions in both scenarios, we have

$$\begin{aligned}
\text{Complete Scenarios: } \mathbb{E}_i[U_{ij}|X_j] &= \sum_{x \in A \setminus Q} (\beta_i^x + \beta_i^Q \gamma_i^x) x_j + \mathbb{E}_i(\epsilon_{ij}|X_j) \\
\text{Incomplete Scenarios: } \mathbb{E}_i[U_{ij}|\tilde{X}_j] &= \sum_{x \in A \setminus Q} (\beta_i^x + \beta_i^Q \tilde{\gamma}_i^x) x_j + \mathbb{E}_i(\epsilon_{ij}|\tilde{X}_j)
\end{aligned} \tag{26}$$

Then, with the same set of identifying assumptions, given the reported compensating differentials and normalizing $\beta^W = 1$, we have the following indifference conditions:

$$\begin{aligned}
\Delta_{ijk} &= \sum_{x \in A \setminus Q} \underbrace{(\beta_i^x + \beta_i^Q \gamma_i^x)}_{\beta_i^{x(CS)}} (x_j - x_k) \\
\tilde{\Delta}_{ijk} &= \sum_{x \in A \setminus Q} \underbrace{(\beta_i^x + \beta_i^Q \tilde{\gamma}_i^x)}_{\beta_i^{x(IS)}} (x_j - x_k)
\end{aligned} \tag{27}$$

The differences in the coefficients in front of the gender differences across the complete and incomplete scenarios give us

$$\beta_i^{G(CS)} - \beta_i^{G(IS)} = \beta_i^Q (\gamma_i^G - \tilde{\gamma}_i^G) \tag{28}$$

γ encapsulates the information on manager quality given X_j , whereas $\tilde{\gamma}$ encapsulates the information on manager quality given \tilde{X}_j , i.e., in the absence of the information on manager quality. In the presence of statistical discrimination against female managers, this should be negative. This is what our estimates show, given that individuals care positively about manager quality. Thus, under this model specification, statistical discrimination is identified. Observe the analogy with the model presented in the main paper. Here, too, if the individual does not care about manager quality (i.e., $\beta^Q = 0$), the parameters identified from the complete and incomplete scenarios must be identical because the variation in information revelation will have no effect.

A.5 Details on survey administration and data collection

The online survey was designed and implemented on Qualtrics. Recruitment of students was done by the research assistants (RAs), who had previous experience in recruiting students for surveys and RCTs. In the midst of a nationwide lockdown, we administered the online survey in three key steps.

Step 1: The RAs, based on their previous experience, sent a sign-up link to each department's student class representative, who distributed the link in the class lists. The sign-up sheet, in addition to containing the consent form for them to sign, asked for e-mail addresses to which the link of the survey would be sent, as well as basic demographic information, department affiliation, faculty of study (arts/science/engineering) and level (bachelor's or master's) and year of study. Students were also allowed to choose the date and the time at which they would like to take the survey. They had a choice among 4 dates from April 8th to April 11th. On their selected date, they had a choice among 6 time slots: 10 am, 12 noon, 5 pm, 7 pm, 9 pm and 11 pm. We observed that our pilots that had specified time slots along with dates had higher completion rates than those with just dates. The sign-up form ended with a summary. The sign-up form was designed to automatically send respondents' enrollment form to their email address. This enabled us to automatically have the signed copy of the consent form sent to the participant.

Step 2: Upon receiving the sign-ups, we scheduled emails to be sent out with unique links to the survey for each participant an hour before each he or she was scheduled to participate in the survey. Hence, the link could not be used on two different devices to fill in the survey. The survey was also designed to prevent ballot-boxing; i.e., once the survey was completed from one link, when clicked again, that link would show a confirmation that the survey had already been completed. The links were designed to expire within 24 hours. Thanks to the extensive pilots done before, we did not face any technical difficulties while implementing the survey. Debriefs with pilot participants were extremely helpful for rewording the questions to optimize communication and maximize participants' understanding.

Step 3: The mode and details of online payment were selected in the last section of the survey. The options included direct bank transfers, PayPal and UPI (unified payment interface). The payment was processed for the list of verified students within the prestated timeline for each payment mode.